



LFS258

Kubernetes Fundamentals

Version 2018-11-13



© Copyright the Linux Foundation 2018. All rights reserved.

The training materials provided or developed by The Linux Foundation in connection with the training services are protected by copyright and other intellectual property rights.

Open source code incorporated herein may have other copyright holders and is used pursuant to the applicable open source license.

The training materials are provided for individual use by participants in the form in which they are provided. They may not be copied, modified, distributed to non-participants or used to provide training to others without the prior written consent of The Linux Foundation.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the **Linux Foundation**

<http://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please contact:

training@linuxfoundation.org

Contents

- 1 Introduction 1**
 - 1.1 Labs 1
- 2 Basics of Kubernetes 3**
 - 2.1 Labs 3
- 3 Installation and Configuration 5**
 - 3.1 Labs 5
- 4 Kubernetes Architecture 19**
 - 4.1 Labs 19
- 5 APIs and Access 31**
 - 5.1 Labs 31
- 6 API Objects 37**
 - 6.1 Labs 37
- 7 Managing State With Deployments 45**
 - 7.1 Labs 45
- 8 Services 53**
 - 8.1 Labs 53
- 9 Volumes and Data 59**
 - 9.1 Labs 59
- 10 Ingress 75**
 - 10.1 Labs 75
- 11 Scheduling 81**
 - 11.1 Labs 81
- 12 Logging and Troubleshooting 89**
 - 12.1 Labs 89
- 13 Custom Resource Definition 93**
 - 13.1 Labs 93
- 14 Kubernetes Federation 97**

14.1 Labs	97
15 Helm	99
15.1 Labs	99
16 Security	105
16.1 Labs	105

List of Figures

3.1 External Access via Browser 18

Chapter 1

Introduction



1.1 Labs

Exercise 1.1: Configuring the System for **sudo**

It is very dangerous to run a **root shell** unless absolutely necessary: a single typo or other mistake can cause serious (even fatal) damage.

Thus, the sensible procedure is to configure things such that single commands may be run with superuser privilege, by using the **sudo** mechanism. With **sudo** the user only needs to know their own password and never needs to know the root password.

If you are using a distribution such as **Ubuntu**, you may not need to do this lab to get **sudo** configured properly for the course. However, you should still make sure you understand the procedure.

To check if your system is already configured to let the user account you are using run **sudo**, just do a simple command like:

```
$ sudo ls
```

You should be prompted for your user password and then the command should execute. If instead, you get an error message you need to execute the following procedure.

Launch a root shell by typing **su** and then giving the **root** password, not your user password.

On all recent **Linux** distributions you should navigate to the `/etc/sudoers.d` subdirectory and create a file, usually with the name of the user to whom root wishes to grant **sudo** access. However, this convention is not actually necessary as **sudo** will scan all files in this directory as needed. The file can simply contain:

```
student ALL=(ALL) ALL
```

if the user is student.

An older practice (which certainly still works) is to add such a line at the end of the file `/etc/sudoers`. It is best to do so using the **visudo** program, which is careful about making sure you use the right syntax in your edit.

You probably also need to set proper permissions on the file by typing:

```
$ chmod 440 /etc/sudoers.d/student
```

(Note some **Linux** distributions may require 400 instead of 440 for the permissions.)

After you have done these steps, exit the root shell by typing `exit` and then try to do `sudo ls` again.

There are many other ways an administrator can configure **sudo**, including specifying only certain permissions for certain users, limiting searched paths etc. The `/etc/sudoers` file is very well self-documented.

However, there is one more setting we highly recommend you do, even if your system already has **sudo** configured. Most distributions establish a different path for finding executables for normal users as compared to root users. In particular the directories `/sbin` and `/usr/sbin` are not searched, since **sudo** inherits the `PATH` of the user, not the full root user.

Thus, in this course we would have to be constantly reminding you of the full path to many system administration utilities; any enhancement to security is probably not worth the extra typing and figuring out which directories these programs are in. Consequently, we suggest you add the following line to the `.bashrc` file in your home directory:

```
PATH=$PATH:/usr/sbin:/sbin
```

If you log out and then log in again (you don't have to reboot) this will be fully effective.

Chapter 2

Basics of Kubernetes



2.1 Labs

Exercise 2.1: View Online Resources

Visit kubernetes.io

With such a fast changing project, it is important to keep track of updates. The main place to find documentation of the current version is <https://kubernetes.io/>.

1. Open a browser and visit the <https://kubernetes.io/> website.
2. In the upper right hand corner, use the drop down to view the versions available. It will say something like v1.12.
3. Select the top level link for Documentation. The links on the left of the page can be helpful in navigation.
4. As time permits navigate around other sub-pages such as SETUP, CONCEPTS, and TASKS to become familiar with the layout.

Track Kubernetes Issues

There are hundreds, perhaps thousands, working on Kubernetes every day. With that many people working in parallel there are good resources to see if others are experiencing a similar outage. Both the source code as well as feature and issue tracking are currently on [github.com](https://github.com/kubernetes/kubernetes).

1. To view the main page use your browser to visit <https://github.com/kubernetes/kubernetes/>
2. Click on various sub-directories and view the basic information available.
3. Update your URL to point to <https://github.com/kubernetes/kubernetes/issues>. You should see a series of issues, feature requests, and support communication.

4. In the search box you probably see some existing text like `is:issue is:open` which allows you to filter on the kind of information you would like to see. Append the search string to read: `is:issue is:open label:kind/bug` then press enter.
5. You should now see bugs in descending date order. Across the top of the issues a menu area allows you to view entries by author, labels, projects, milestones, and assignee as well. Take a moment to view the various other selection criteria.
6. Some times you may want to exclude a kind of output. Update the URL again, but precede the label with a minus sign, like: `is:issue is:open -label:kind/bug`. Now you see everything except bug reports.
7. Explore the page with the remaining time left.

Chapter 3

Installation and Configuration



3.1 Labs

Exercise 3.1: Install Kubernetes

Overview

There are several Kubernetes installation tools provided by various vendors. In this lab we will learn to use **kubeadm**. As a community-supported independent tool, it is planned to become the primary manner to build a Kubernetes cluster.

The labs were written using **Ubuntu** instances running on **Google Cloud Platform (GCP)**. They have been written to be vendor-agnostic so could run on AWS, local hardware, or inside of virtualization to give you the most flexibility and options. Each platform will have different access methods and considerations. As of v1.12.1 the minimum (as in barely works) size for **VirtualBox** is 3vCPU/1G memory/5G minimal OS for master and 1vCPU/1G memory/5G minimal OS for worker node.

If using your own equipment you will have to disable swap on every node. There may be other requirements which will be shown as warnings or errors when using the **kubeadm** command. While most commands are run as a regular user, there are some which require root privilege. Please configure **sudo** access as shown in a previous lab. You If you are accessing the nodes remotely, such as with **GCP** or **AWS**, you will need to use an SSH client such as a local terminal or **PuTTY** if not using **Linux** or a Mac. You can download **PuTTY** from www.putty.org. You would also require a **.pem** or **.ppk** file to access the nodes. Each cloud provider will have a process to download or create this file. If attending in-person instructor led training the file will be made available during class.

In the following exercise we will install Kubernetes on a single node then grow the cluster, adding more compute resources. Both nodes used are the same size, providing 2 vCPUs and 7.5G of memory. Smaller nodes could be used, but would run slower.

Various exercises will use YAML files, which are included in the text. You are encouraged to write the files when possible, as the syntax of YAML has white space indentation requirements that are important to learn. An important note, **do not** use tabs in your YAML files, **white space only. Indentation matters.**

If using a PDF the use of copy and paste often does not paste the single quote correctly. It pastes as a back-quote instead. You will need to modify it by hand. The files have also been made available as a compressed **tar** file. You can view the resources by navigating to this URL:

<https://training.linuxfoundation.org/cm/LFS258>

To login use user: LFtraining and a password of: Penguin2014

Once you find the name and link of the current file, which will change as the course updates, use **wget** to download the file into your node from the command line then expand it like this:

```
$ wget https://training.linuxfoundation.org/cm/LFS258/LFS258.V2018-11-13_SOLUTIONS.tar.bz2 \
    --user=LFtraining --password=Penguin2014

$ tar -xvf LFS258.V2018-11-13_SOLUTIONS.tar.bz2
```

(Note: depending on your software, if you are cutting and pasting the above instructions, the underscores may disappear and be replaced by spaces, so you may have to edit the command line by hand!)

While **Ubuntu 18** bionic has become the typical version to deploy the Kubernetes repository does not yet have compatible binaries at the time of this writing. While **xenial** binaries can be used there are many additional steps necessary to complete the labs. **Ubuntu 18** is expected to be available by the time Kubernetes v.1.13 is released.

Install Kubernetes

Log into your nodes. If attending in-person instructor led training the node IP addresses will be provided by the instructor. You will need to use a **.pem** or **.ppk** key for access, depending on if you are using **ssh** from a terminal or **PuTTY**. The instructor will provide this to you.

1. Open a terminal session on your first node. For example, connect via **PuTTY** or **SSH** session to the first **GCP** node. The user name may be different than the one shown, **student**. The IP used in the example will be different than the one you will use.

```
[student@laptop ~]$ ssh -i LFS458.pem student@35.226.100.87
The authenticity of host '54.214.214.156 (35.226.100.87)' can't be established.
ECDSA key fingerprint is SHA256:IPvznbkx93/Wc+ACwXrCcDDgvBwmvEXC9vmYhk2Wo1E.
ECDSA key fingerprint is MD5:d8:c9:4b:b0:b0:82:d3:95:08:08:4a:74:1b:f6:e1:9f.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '35.226.100.87' (ECDSA) to the list of known hosts.
<output_omitted>
```

2. Become root and update and upgrade the system. Answer any questions to use the defaults.

```
student@lfs458-node-1a0a:~$ sudo -i

root@lfs458-node-1a0a:~# apt-get update && apt-get upgrade -y
<output_omitted>
```

3. The main choices for a container environment are **Docker** and **cri-o**. We will use **Docker** for class, as **cri-o** requires a fair amount of extra work to enable for Kubernetes. As **cri-o** is open source the community seems to be heading towards its use.

```
root@lfs458-node-1a0a:~# apt-get install -y docker.io
<output-omitted>
```

4. Add new repo for kubernetes. You could also get a tar file or use code from GitHub. Create the file and add an entry for the main repo for your distribution. As we are still using Ubuntu 16.04 add the kubernetes-xenial with the key word main. Note there are four sections to the entry.

```
root@lfs458-node-1a0a:~# vim /etc/apt/sources.list.d/kubernetes.list
deb http://apt.kubernetes.io/ kubernetes-xenial main
```

5. Add a GPG key for the packages. The command spans three lines. You can omit the backslash when you type. The OK is the expected output, not part of the command.

```
root@lfs458-node-1a0a:~# curl -s \
    https://packages.cloud.google.com/apt/doc/apt-key.gpg \
    | apt-key add -
OK
```

6. Update with new repo, which will download new repo information.

```
root@lfs458-node-1a0a:~# apt-get update
<output-omitted>
```

7. Install the software. There are regular releases the newest of which can be used by omitting the equal sign and version information on the command line. Historically new version have lots of changes and a good chance of a bug or five.

```
root@lfs458-node-1a0a:~# apt-get install -y \
    kubeadm=1.12.1-00 kubelet=1.12.1-00 kubectl=1.12.1-00
<output-omitted>
```

8. Deciding which pod network to use for Container Networking Interface (CNI) should take into account the expected demands on the cluster. There can be only one pod network per cluster, although the **CNI-Genie** project is trying to change this.

The network must allow container-to-container, pod-to-pod, pod-to-service, and external-to-service communications. As **Docker** uses host-private networking, using the `docker0` virtual bridge and `veth` interfaces would require being on that host to communicate.

We will use **Calico** as a network plugin which will allow us to use Network Policies later in the course. Currently **Calico** does not deploy using CNI by default. The 3.3 version of **Calico** has more than one configuration file for flexibility with RBAC. Download the configuration files for. Once downloaded look for the expected IPV4 range for containers to use.

A short url for each file is shown, the longer URLs can be found here: <https://docs.projectcalico.org/v3.3/getting-started/kubernetes/installation/hosted/rbac-kdd.yaml> and: <https://docs.projectcalico.org/v3.3/getting-started/kubernetes/installation/hosted/kubernetes-datastore/calico-networking/1.7/calico.yaml>

```
root@lfs458-node-1a0a:~# wget https://tinyurl.com/yb4xturm \
    -O rbac-kdd.yaml
```

```
root@lfs459-node-1a0a:~# wget https://tinyurl.com/y8lvqc9g \
    -O calico.yaml}
```

9. Use **less** to page through the file. Look for the IPV4 pool expected by the containers. There are many different configuration settings in this file. Take a moment to view the entire file. The `CALICO_IPV4POOL_CIDR` must match the value given to **kubeadm init** in the following step, whatever the value may be.

```
root@lfs458-node-1a0a:~# less calico.yaml
....
    # Configure the IP Pool from which Pod IPs will be chosen.
    - name: CALICO_IPV4POOL_CIDR
      value: "192.168.0.0/16"
....
```

10. Initialize the master. Read through the output line by line. Expect the output to change as the software matures. At the end are configuration directions to run as a non-root user. The token is mentioned as well. This information can be found later with the **kubeadm token list** command. The output also directs you to create a pod network to the cluster, which will be our next step. Pass the network settings **Calico** has in its configuration file, found in the previous step. **Please note:** the output lists several commands which following commands will complete. Read the next step before further typing.

```
root@lfs458-node-1a0a:~# kubeadm init --pod-network-cidr 192.168.0.0/16
```

```
[init] using Kubernetes version: v1.12.1
[preflight] running pre-flight checks
[preflight/images] Pulling images required for setting up a
  Kubernetes cluster
[preflight/images] This might take a minute or two, depending
  on the speed of your internet connection
```

<output-omitted>

Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of machines by running the following on each node as root:

```
kubeadm join 10.128.0.4:6443 --token rdnhok.g8mb6lfgesunanvh
--discovery-token-ca-cert-hash
sha256:66350d154fc0169b5bb5fd50c04b72468195e356d78d95f137ed55e995402f77
```

11. As suggested in the directions at the end of the previous output we will allow a non-root user admin level access to the cluster. Take a quick look at the configuration file once it has been copied and the permissions fixed.

```
root@lfs458-node-1a0a:~# exit
logout
```

```
student@lfs458-node-1a0a:~$ mkdir -p $HOME/.kube
```

```
student@lfs458-node-1a0a:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
student@lfs458-node-1a0a:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
student@lfs458-node-1a0a:~$ less .kube/config
apiVersion: v1
clusters:
- cluster:
<output_omitted>
```

12. Apply the network plugin configuration to your cluster. Remember to copy the file to the current, non-root user directory first.

```
student@lfs458-node-1a0a:~$ sudo cp /root/rbac-kdd.yaml .
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f rbac-kdd.yaml
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
```

```
student@lfs458-node-1a0a:~$ sudo cp /root/calico.yaml .
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f calico.yaml
configmap/calico-config created
service/calico-typha created
deployment.apps/calico-typha created
poddisruptionbudget.policy/calico-typha created
<output_omitted>
```

- While many objects have short names, a **kubectl** command can be a lot to type. We will enable **bash** auto-completion. Begin by adding the settings to the current shell. Then update the `~/.bashrc` file to make it persistent.

```
student@lfs458-node-1a0a:~$ source <(kubectl completion bash)

student@lfs458-node-1a0a:~$ echo "source <(kubectl completion bash)" >> ~/.bashrc
```

- Test by describing the node again. Type the first three letters of the sub-command then type the **Tab** key. Auto-completion assumes the default namespace. Pass the namespace first to use auto-completion with a different namespace. By pressing **Tab** multiple times you will see a list of possible values. Continue typing until a unique name is used.

```
student@lfs458-node-1a0a:~$ kubectl des<Tab> n<Tab><Tab> lfs458-<Tab>

student@lfs458-node-1a0a:~$ kubectl -n kube-s<Tab> g<Tab> po e<Tab>
```

Exercise 3.2: Grow the Cluster

Open another terminal and connect into a your second node. Install **Docker** and Kubernetes software. These are the many, but not all, of the steps we did on the master node.

The book will use the **lfs458-worker** prompt for the node being added to help keep track of the proper node for each command. Note that the prompt indicates both the user and system upon which run the command.

- Using the same process as before connect to a second node. If attending ILT use the same `.pem` key and a new IP provided by the instructor to access the new node. Giving a title or color to the new terminal window is probably a good idea to keep track of the two systems. The prompts can look very similar.

```
student@lfs458-worker:~$ sudo -i

root@lfs458-worker:~# apt-get update && apt-get upgrade -y

root@lfs458-worker:~# apt-get install -y docker.io

root@lfs458-worker:~# vim /etc/apt/sources.list.d/kubernetes.list
deb http://apt.kubernetes.io/ kubernetes-xenial main

root@lfs458-worker:~# curl -s \
    https://packages.cloud.google.com/apt/doc/apt-key.gpg \
    | apt-key add -

root@lfs458-worker:~# apt-get update

root@lfs458-worker:~# apt-get install -y \
    kubeadm=1.12.1-00 kubelet=1.12.1-00 kubectl=1.12.1-00
```

- Find the IP address of your master server. The interface name will be different depending on where the node is running. Currently inside of **GCE** the primary interface for this node type is `ens4`. Your interfaces names may be different. From the output we know our master node IP is `10.128.0.3`.

```
student@lfs458-node-1a0a:~$ ip addr show ens4 | grep inet
    inet 10.128.0.3/32 brd 10.128.0.3 scope global ens4
    inet6 fe80::4001:aff:fe8e:2/64 scope link
```

- Find the token on the master node. The token lasts 24 hours by default. If it has been longer, and no token is present you can generate a new one with the **sudo kubeadm token create** command, seen in the following command.

```
student@lfs458-node-1a0a:~$ sudo kubeadm token list
TOKEN                                TTL      EXPIRES      USAGES      DESCRIPTION
27eee4.6e66ff60318da929             23h      2017-11-03T13:27:33Z
authentication,signing The default bootstrap token generated
by 'kubeadm init'....
```

4. **Only if the token has expired**, you can create a new token, to use as part of the join command.

```
student@lfs458-node-1a0a:~$ sudo kubeadm token create
27eee4.6e66ff60318da929
```

5. Starting in v1.9 you should create and use a Discovery Token CA Cert Hash created from the master to ensure the node joins the cluster in a secure manner. Run this on the master node or wherever you have a copy of the CA file. You will get a long string as output.

```
student@lfs458-node-1a0a:~$ openssl x509 -pubkey \
    -in /etc/kubernetes/pki/ca.crt | openssl rsa \
    -pubin -outform der 2>/dev/null | openssl dgst \
    -sha256 -hex | sed 's/^.* //'
6d541678b05652e1fa5d43908e75e67376e994c3483d6683f2a18673e5d2a1b0
```

6. Use the token and hash, in this case as sha256:<hash> to join the cluster from the second node. Use the **private** IP address of the master server and port 6443. The output of the **kubeadm init** on the master also has an example to use, should it still be available.

```
root@lfs458-worker:~# kubeadm join \
    --token 27eee4.6e66ff60318da929 \
    10.128.0.3:6443 \
    --discovery-token-ca-cert-hash \
    sha256:6d541678b05652e1fa5d43908e75e67376e994c3483d6683f2a18673e5d2a1b0
```

```
[preflight] Running pre-flight checks.
[WARNING FileExisting-crictl]: crictl not found in system path
[discovery] Trying to connect to API Server "10.142.0.2:6443"
[discovery] Created cluster-info discovery client, requesting info from
"https://10.142.0.2:6443"
[discovery] Requesting info from "https://10.142.0.2:6443" again to
validate TLS against the pinned public key
[discovery] Cluster info signature and contents are valid and TLS
certificate validates against pinned roots, will
use API Server "10.142.0.2:6443"
[discovery] Successfully established connection with API Server
"10.142.0.2:6443"
This node has joined the cluster:
* Certificate signing request was sent to master and a response
  was received.
* The Kubelet was informed of the new secure connection details.
Run 'kubectl get nodes' on the master to see this node join the cluster.
```

7. Try to run the **kubectl** command on the secondary system. It should fail. You do not have the cluster or authentication keys in your local **.kube/config** file.

```
root@lfs458-worker:~# exit

student@lfs458-worker:~$ kubectl get nodes
The connection to the server localhost:8080 was refused
- did you specify the right host or port?

student@lfs458-worker:~$ ls -l .kube
ls: cannot access '.kube': No such file or directory
```

Exercise 3.3: Finish Cluster Setup

1. View the available nodes of the cluster. It can take a minute or two for the status to change from NotReady to Ready. The NAME field can be used to look at the details. Your node name will be different. Note the master node says NotReady, which is due to a taint.


```
student@lfs458-node-1a0a:~$ kubectl get node
NAME                STATUS    ROLES    AGE   VERSION
lfs458-node-1a0a    NotReady  master   18m   v1.12.1
lfs458-worker       Ready     <none>    3m25s v1.12.1
```

- Look at the details of the node. Work line by line to view the resources and their current status. Notice the status of Taints. The master won't allow non-internal pods by default for security reasons. Take a moment to read each line of output, some appear to be an error until you notice the status shows False.

```
student@lfs458-node-1a0a:~$ kubectl describe node lfs458-node-1a0a
Name:                lfs458-node-1a0a
Roles:               master
Labels:              beta.kubernetes.io/arch=amd64
                    beta.kubernetes.io/os=linux
                    kubernetes.io/hostname=lfs458-node-1a0a
                    node-role.kubernetes.io/master=
Annotations:         kubeadm.alpha.kubernetes.io/cri-socket=/var/run/dockershim.sock
                    node.alpha.kubernetes.io/ttl=0
                    volumes.kubernetes.io/controller-managed-attach-detach=true
CreationTimestamp:   Sun, 29 Jul 2018 21:29:32 +0000
Taints:              node-role.kubernetes.io/master:NoSchedule
<output_omitted>
```

- Allow the master server to run non-infrastructure pods. The master node begins tainted for security and performance reasons. Will allow usage of the node in the training environment, but this step may be skipped in a production environment. Note the minus sign (-) at the end, which is the syntax to remove a taint. As the second node does not have the taint you will get a not found error.

```
student@lfs458-node-1a0a:~$ kubectl describe node | grep -i taint
Taints:              node-role.kubernetes.io/master:NoSchedule
Taints:              <none>
```

```
student@lfs458-node-1a0a:~$ kubectl taint nodes \
    --all node-role.kubernetes.io/master-

node/lfs458-node-1a0a untainted
error: taint "node-role.kubernetes.io/master:" not found
```

- Now that the master node is able to execute any pod we **may** find there is a new taint. This behavior began with v1.12.0, requiring a newly added node to be enabled. View then remove the taint if present. It can take a minute or two for the scheduler to deploy the remaining pods.

```
student@lfs458-node-1a0a:~$ kubectl describe node | grep -i taint
Taints:              node.kubernetes.io/not-ready:NoSchedule
Taints:              <none>
```

```
student@lfs458-node-1a0a:~$ kubectl taint nodes \
    --all node.kubernetes.io/not-ready-

node/lfs58-node-1a0a untainted
error: taint "node.kubernetes.io/not-ready:" not found
```

- Another "undocumented feature" in v1.12.1 is that the taint removal does not always work the first time. Check to see if the taint has been removed. You may have to remove the taint two or three times before it is actually gone. Wait 60 seconds before trying again to remove the taint.

```
student@lfs458-node-1a0a:~$ kubectl describe node | grep -i taint
Taints:              node.kubernetes.io/not-ready:NoSchedule
Taints:              <none>
```

```
student@lfs458-node-1a0a:~$ kubectl taint nodes \
--all node.kubernetes.io/not-ready-

node/lfs58-node-1a0a untainted
error: taint "node.kubernetes.io/not-ready:" not found

student@lfs458-node-1a0a:~$ sleep 60 ; kubectl describe node | grep -i taint
Taints:          <none>
Taints:          <none>
```

6. Determine if the DNS and Calico pods are ready for use. They should all show a status of Running. It may take a minute or two to transition from Pending.

```
student@lfs458-node-1a0a:~$ kubectl get pods --all-namespaces
NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
kube-system    calico-etcd-jlgwr                      1/1     Running   0          6m
kube-system    calico-kube-controllers-74b88b647-wlqf5 1/1     Running   0          6m
kube-system    calico-node-tpvnr                      2/2     Running   0          6m
kube-system    coredns-78fcd6894-nc5cn                1/1     Running   0          17m
kube-system    coredns-78fcd6894-xs96m                1/1     Running   0          17m
<output_omitted>
```

7. If you notice the coredns- pods are stuck in ContainerCreating status you may have to delete them, causing new ones to be generated. Delete both pods and check to see they show a Running state. Your pod names will be different.

```
student@lfs458-node-1a0a:~$ kubectl get pods --all-namespaces
NAMESPACE      NAME                                     READY   STATUS
RESTARTS   AGE
kube-system    calico-node-qkvzh                      2/2     Running
0          59m
kube-system    calico-node-vndn7                      2/2     Running
0          12m
kube-system    coredns-576cbf47c7-rn6v4              0/1     ContainerCreating
0          3s
kube-system    coredns-576cbf47c7-vq5dz              0/1     ContainerCreating
0          94m
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl -n kube-system delete \
pod coredns-576cbf47c7-vq5dz coredns-576cbf47c7-rn6v4
pod "coredns-576cbf47c7-vq5dz" deleted
pod "coredns-576cbf47c7-rn6v4" deleted
```

8. When it finished you should see a new tunnel, tunl0, and a cali interface. It may take up to a minute to be created. As you create objects more interfaces will be created.

```
student@lfs458-node-1a0a:~$ ip a
<output_omitted>
4: tunl0@NONE: <NOARP,UP,LOWER_UP> mtu 1440 qdisc noqueue state
UNKNOWN group default qlen 1000
    link/ipip 0.0.0.0 brd 0.0.0.0
    inet 192.168.0.1/32 brd 192.168.0.1 scope global tunl0
        valid_lft forever preferred_lft forever
6: cali0b93ed4661@if4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu
1440 qdisc noqueue state UP group default
    link/ether ee:ee:ee:ee:ee:ee brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet6 fe80::ecee:eeff:feee:eeee/64 scope link
        valid_lft forever preferred_lft forever
```

Exercise 3.4: Deploy A Simple Application

We will test to see if we can deploy a simple application, in this case the **nginx** web server.

1. Create a new deployment, which is an Kubernetes object while will deploy and monitor an application in a container. Verify it is running and the desired number of container matches the available.

```
student@lfs458-node-1a0a:~$ kubectl create deployment nginx --image=nginx
deployment.apps/nginx created
```

```
student@lfs458-node-1a0a:~$ kubectl get deployments
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx         1         1         1            1           6s
```

2. View the details of the deployment. Remember auto-completion will work for sub-commands and resources as well.

```
student@lfs458-node-1a0a:~$ kubectl describe deployment nginx
Name:          nginx
Namespace:     default
CreationTimestamp: Thu, 08 Nov 2018 19:23:00 +0000
Labels:        app=nginx
Annotations:   deployment.kubernetes.io/revision: 1
Selector:      app=nginx
Replicas:      1 desired | 1 updated | 1 total | 1 ava...
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
<output_omitted>
```

3. View the basic steps the cluster took in order to pull and deploy the new application. You should see several lines of output with newer events at the top.

```
student@lfs458-node-1a0a:~$ kubectl get events
<output_omitted>
```

4. You can also view the output in **yaml** format, which could be used to create this deployment again or new deployments. Get the information but change the output to yaml. Note that halfway down there is status information of the current deployment.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx -o yaml
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  creationTimestamp: 2017-09-27T18:21:25Z
<output_omitted>
```

5. Run the command again and redirect the output to a file. Then edit the file. Remove the creationTimestamp, resourceVersion, selfLink, and uid lines. Also remove all the lines including and after status:, which should be somewhere around line 40, if others have already been removed.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx -o yaml > first.yaml

student@lfs458-node-1a0a:~$ vim first.yaml

<Remove the lines mentioned above>
```

6. Delete the existing deployment.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment nginx
deployment.extensions "nginx" deleted
```

7. Create the deployment again this time using the file.

```
student@lfs458-node-1a0a:~$ kubectl create -f first.yaml
deployment.extensions/nginx created
```

8. Look at the yaml output of this iteration and compare it against the first. The time stamp, resource version and uid we had deleted are in the new file. These are generated for each resource we create, so we need to delete them from yaml files to avoid conflicts or false information. The status should not be hard-coded either.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx -o yaml > second.yaml

student@lfs458-node-1a0a:~$ diff first.yaml second.yaml
<output_omitted>
```

9. Now that we have worked with the raw output we will explore two other ways of generating useful YAML or JSON. Use the `--dry-run` option and verify no object was created. Only the prior `nginx` deployment should be found. The output lacks the unique information we removed before.

```
student@lfs458-node-1a0a:~$ kubectl create deployment two --image=nginx --dry-run -o yaml
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    run: two
  name: two
spec:
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl get deployment
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx         1         1         1            1           7m
```

10. Existing objects can be viewed in a ready to use YAML output. Take a look at the existing `nginx` deployment. Note there is more detail to the `--export` option.

```
student@lfs458-node-1a0a:~$ kubectl get deployments nginx --export -o yaml
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  creationTimestamp: null
  generation: 1
  labels:
    run: nginx
<output_omitted>
```

11. The output can also be viewed in JSON output.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx --export -o json
{
  "apiVersion": "extensions/v1beta1",
  "kind": "Deployment",
  "metadata": {
    "annotations": {
      "deployment.kubernetes.io/revision": "1"
    },
    "creationTimestamp": null,
    "generation": 1,
    "labels": {
      "run": "nginx"
    },
    "name": "nginx"
  },
  "spec": {
    "replicas": 1,
    "selector": {
      "matchLabels": {
        "run": "nginx"
      }
    },
    "template": {
      "metadata": {
        "labels": {
          "run": "nginx"
        }
      },
      "spec": {
        "containers": [
          {
            "image": "nginx",
            "name": "nginx"
          }
        ]
      }
    }
  }
}
<output_omitted>
```

12. The newly deployed `nginx` container is a light weight web server. We will need to create a service to view the default welcome page. Begin by looking at the help output. Note that there are several examples given, about halfway through the output.

```
student@lfs458-node-1a0a:~$ kubectl expose -h
<output_omitted>
```

13. Now try to gain access to the web server. As we have not declared a port to use you will receive an error.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment/nginx
error: couldn't find port via --port flag or introspection
See 'kubectl expose -h' for help and examples.
```

14. To change an existing configuration in a cluster can be done with subcommands `apply`, `edit` or `patch` for non-disruptive updates. The `apply` command does a three-way diff of previous, current, and supplied input to determine modifications to make. Fields not mentioned are unaffected. The `edit` function performs a `get`, opens an editor, then an `apply`. You can update API objects in place with JSON `patch` and `merge patch` or `strategic merge patch` functionality.

If the configuration has resource fields which cannot be updated once initialized then a disruptive update could be done using the `replace --force` option. This deletes first then re-creates a resource.

Edit the file. Find the container name, somewhere around line 31 and add the port information as shown below.

```
student@lfs458-node-1a0a:~$ vim first.yaml
....
spec:
  containers:
  - image: nginx
    imagePullPolicy: Always
    name: nginx
    ports:                                # Add these
    - containerPort: 80                  # three
      protocol: TCP                      # lines
    resources: {}
....
```

15. Due to how the object was created we will need to use `replace` to terminate and create a new deployment.

```
student@lfs458-node-1a0a:~$ kubectl replace -f first.yaml
deployment.extensions/nginx replaced
```

16. View the Pod and Deployment. Note the AGE shows the Pod was re-created.

```
student@lfs458-node-1a0a:~$ kubectl get deploy,pod
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/nginx	1	1	1	1	2m4s

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-7cbc4b4d9c-l8cgl	1/1	Running	0	8s

17. Try to expose the resource again. This time it should work.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment/nginx
service/nginx exposed
```

18. Verify the service configuration. First look at the service information, then at the endpoint information. Note the Cluster IP is not the current endpoint. Take note of the current endpoint IP. In the example below it is `10.244.1.99:80`. We will use this information in a few steps.

```
student@lfs458-node-1a0a:~$ kubectl get svc nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx	ClusterIP	10.100.61.122	<none>	80/TCP	3m


```
student@lfs458-node-1a0a:~$ kubectl get ep nginx
```

NAME	ENDPOINTS	AGE
nginx	10.244.1.99:80	4m

19. Determine which node the container is running on. Log into that node and use `tcpdump` to view traffic on the `tun10`, as in tunnel zero, interface. The second node in this example. You may also see traffic on an interface which starts with `cali` and some string. Leave that command running while you run `curl` in the following step. You should see several messages go back and forth, including a `HTTP: HTTP/1.1 200 OK` and a `ack` response to the same sequence.

```
student@lfs458-node-1a0a:~$ kubectl describe pod nginx-7cbc4b4d9c-d27xw \
| grep Node:
Node:    lfs458-worker/10.128.0.5

student@lfs458-worker:~$ sudo tcpdump -i tunl0
tcpdump: verbose output suppressed, use -v or -vv for full protocol...
listening on tunl0, link-type EN10MB (Ethernet), capture size...
<output_omitted>
```

20. Test access to the Cluster IP, port 80. You should see the generic `nginx` installed and working page. The output should be the same when you look at the `ENDPOINTS` IP address. If the `curl` command times out the pod may be running on the other node. Run the same command on that node and it should work.

```
student@lfs458-node-1a0a:~$ curl 10.100.61.122:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>
```

21. Now scale up the deployment from one to three web servers.

```
student@lfs458-node-1a0a:~$ kubectl get deployment nginx
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx         1         1         1            1           12m

student@lfs458-node-1a0a:~$ kubectl scale deployment nginx --replicas=3
deployment.extensions/nginx scaled

student@lfs458-node-1a0a:~$ kubectl get deployment nginx
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx         3         3         3            3           12m
```

22. View the current endpoints. There now should be three. If the `DESIRED` above said three, but `AVAILABLE` said two wait a few seconds and try again, it could be slow to fully deploy.

```
student@lfs458-node-1a0a:~$ kubectl get ep nginx
NAME          ENDPOINTS                                     AGE
nginx         10.244.0.66:80,10.244.1.100:80,10.244.1.99:80 10m
```

23. Find the oldest pod of the `nginx` deployment and delete it. The `Tab` key can be helpful for the long names. Use the `AGE` field to determine which was running the longest. You will notice activity in the other terminal where `tcpdump` is running, when you delete the pod.

```
student@lfs458-node-1a0a:~$ kubectl get po -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP
nginx-1423793266-7f1qw  1/1    Running   0          14m   10.244.0.66
nginx-1423793266-8w2nk  1/1    Running   0          1m    10.244.1.100
nginx-1423793266-fbt4b  1/1    Running   0          1m    10.244.1.101

student@lfs458-node-1a0a:~$ kubectl delete po nginx-1423793266-7f1qw
pod "nginx-1423793266-7f1qw" deleted
```

24. Wait a minute or two then view the pods again. One should be newer than the others. In the following example two minutes instead of four. If your `tcpdump` was using the `veth` interface of that container it will error out.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY   STATUS    RESTARTS   AGE
nginx-1423793266-13p69  1/1    Running   0          2m
nginx-1423793266-8w2nk  1/1    Running   0          4m
nginx-1423793266-fbt4b  1/1    Running   0          4m
```

25. View the endpoints again. The original endpoint IP is no longer in use. You can delete any of the pods and the service will forward traffic to the existing backend pods.

```
student@lfs458-node-1a0a:~$ kubectl get ep nginx
NAME           ENDPOINTS                                     AGE
nginx          10.244.0.66:80,10.244.1.100:80,10.244.1.101:80 15m
```

26. Test access to the web server again, using the ClusterIP address, then any of the endpoint IP addresses. Even though the endpoints have changed you still have access to the web server. This access is only from within the cluster.

```
student@lfs458-node-1a0a:~$ curl 10.100.61.122:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
<output_omitted>
```

Exercise 3.5: Access from Outside the Cluster

You can access a Service from outside the cluster using a DNS add-on or **vi** environment variables. We will use environment variables to gain access to a Pod.

1. Begin by getting a list of the pods.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nginx-1423793266-13p69             1/1     Running   0           8m
nginx-1423793266-8w2nk             1/1     Running   0           10m
nginx-1423793266-fbt4b             1/1     Running   0           10m
```

2. Choose one of the pods and use the `exec` command to run **printenv** inside the pod. The following example uses the first pod listed above.

```
student@lfs458-node-1a0a:~$ kubectl exec nginx-1423793266-13p69 \
-- printenv |grep KUBERNETES
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_SERVICE_HOST=10.96.0.1
KUBERNETES_SERVICE_PORT=443
NGINX_SERVICE_HOST=10.100.61.122
NGINX_SERVICE_PORT=80
<output_omitted>
```

3. Find and then delete the existing service for **nginx**.

```
student@lfs458-node-1a0a:~$ kubectl get svc
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kubernetes ClusterIP   10.96.0.1    <none>        443/TCP    2d
nginx     ClusterIP   10.100.61.122 <none>        80/TCP     40m
```

4. Delete the service.

```
student@lfs458-node-1a0a:~$ kubectl delete svc nginx
service "nginx" deleted
```

5. Create the service again, but this time pass the `LoadBalancer` type. Check to see the status and note the external ports mentioned. The output will show the `External-IP` as `pending`. Unless a provider responds with a load balancer it will continue to show as `pending`.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment nginx --type=LoadBalancer
service/nginx exposed
```

```
student@lfs458-node-1a0a:~$ kubectl get svc
NAME            TYPE           CLUSTER-IP    EXTERNAL-IP    PORT(S)        AGE
kubernetes      ClusterIP      10.96.0.1     <none>         443/TCP        2d
nginx           LoadBalancer  10.104.249.102 <pending>     80:32753/TCP   2s
```

- Open a browser on your local system, not the GCE node, and use the public IP of your node and port 32753, shown in the output above. If running the labs on a remote system like **AWS** or **GCE** the CLUSTER-IPs are internal. Use the public IP you used with SSH to gain access.

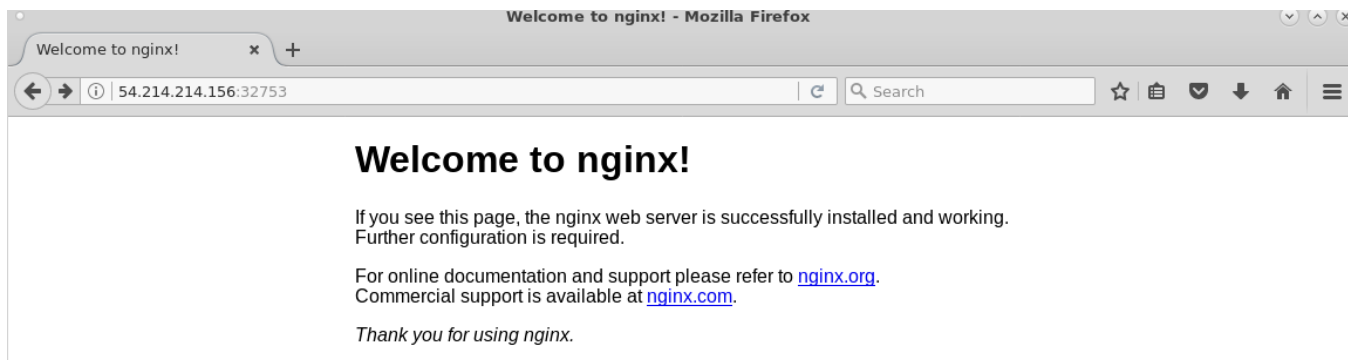


Figure 3.1: **External Access via Browser**

- Scale the deployment to zero replicas. Then test the web page again. It should fail.

```
student@lfs458-node-1a0a:~$ kubectl scale deployment nginx --replicas=0
deployment.extensions/nginx scaled
```

```
student@lfs458-node-1a0a:~$ kubectl get po
No resources found.
```

- Scale the deployment up to two replicas. The web page should work again.

```
student@lfs458-node-1a0a:~$ kubectl scale deployment nginx --replicas=2
deployment.extensions/nginx scaled
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nginx-1423793266-7x181             1/1     Running   0           1m
nginx-1423793266-s6vcz             1/1     Running   0           1m
```

- Delete the deployment to recover system resources. Note that deleting a deployment does not delete the endpoints or services.

```
student@lfs458-node-1a0a:~$ kubectl delete deployments nginx
deployment.extensions "nginx" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl delete ep nginx
endpoints "nginx" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl delete svc nginx
service "nginx" deleted
```


Chapter 4

Kubernetes Architecture



4.1 Labs

Exercise 4.1: Working with CPU and Memory Constraints

Overview

We will continue working with our cluster, which we built in the previous lab. We will work with `resource limits`, more with namespaces and then a complex deployment which you can explore to further understand the architecture and relationships.

Use **SSH** or **PuTTY** to connect to the nodes you installed in the previous exercise. We will deploy an application called **stress** inside a container, and then use `resource limits` to constrain the resources the application has access to use.

1. Use a container called `stress`, which we will name `hog`, to generate load. Verify you have a container running.

```
student@lfs458-node-1a0a:~$ kubectl create deployment hog --image vish/stress
deployment.apps/hog created
```

```
student@lfs458-node-1a0a:~$ kubectl get deployments
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hog           1         1         1            1           12s
```

2. Use the `describe` argument to view details, then view the output in YAML format. Note there are no settings limiting resource usage. Instead, there are empty curly brackets.

```
student@lfs458-node-1a0a:~$ kubectl describe deployment hog
Name:          hog
Namespace:     default
CreationTimestamp: Fri, 09 Nov 2018 19:55:45 +0000
Labels:        app=hog
Annotations:    deployment.kubernetes.io/revision: 1
<output_omitted>
```

```
student@lfs458-node-1a0a:~$ kubectl get deployment hog -o yaml
apiVersion: extensions/v1beta1
```

```

kind: Deployment
Metadata:

<output_omitted>

template:
  metadata:
    creationTimestamp: null
    labels:
      app: hog
  spec:
    containers:
      - image: vish/stress
        imagePullPolicy: Always
        name: stress
        resources: {}
        terminationMessagePath: /dev/termination-log
<output_omitted>

```

3. We will use the YAML output to create our own configuration file. The `--export` option can be useful to not include unique parameters.

```

student@lfs458-node-1a0a:~$ kubectl get deployment hog \
--export -o yaml > hog.yaml

```

4. If you did not use the `--export` option we will need to remove the status output, `creationTimestamp` and other settings, as we don't want to set unique generated parameters. We will also add in memory limits found below.

```

student@lfs458-node-1a0a:~$ vim hog.yaml
.
  imagePullPolicy: Always
  name: hog
  resources:                # Edit to remove {}
    limits:                 # Add these 4 lines
      memory: "4Gi"
    requests:
      memory: "2500Mi"
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
....

```

5. Replace the deployment using the newly edited file.

```

student@lfs458-node-1a0a:~$ kubectl replace -f hog.yaml
deployment.extensions/hog replaced

```

6. Verify the change has been made. The deployment should now show resource limits.

```

student@lfs458-node-1a0a:~$ kubectl get deployment hog -o yaml |less
....
  resources:
    limits:
      memory: 4Gi
    requests:
      memory: 2500Mi
  terminationMessagePath: /dev/termination-log
....

```

7. View the `stdio` of the hog container. Note how much memory has been allocated.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY    STATUS    RESTARTS   AGE
hog-64cbfcc7cf-lwq66 1/1      Running   0           2m

student@lfs458-node-1a0a:~$ kubectl logs hog-64cbfcc7cf-lwq66
I1102 16:16:42.638972      1 main.go:26] Allocating "0" memory, in
    "4Ki" chunks, with a 1ms sleep between allocations
I1102 16:16:42.639064      1 main.go:29] Allocated "0" memory
```

8. Open a second and third terminal to access both master and second nodes. Run **top** to view resource usage. You should not see unusual resource usage at this point. The **dockerd** and **top** processes should be using about the same amount of resources. The **stress** command should not be using enough resources to show up.
9. Edit the hog configuration file and add arguments for **stress** to consume CPU and memory.

```
student@lfs458-node-1a0a:~$ vim hog.yaml
resources:
  limits:
    cpu: "1"
    memory: "4Gi"
  requests:
    cpu: "0.5"
    memory: "500Mi"
args:
- -cpus
- "2"
- -mem-total
- "950Mi"
- -mem-alloc-size
- "100Mi"
- -mem-alloc-sleep
- "1s"
```

10. Delete and recreate the deployment. You should see CPU usage almost immediately and memory allocation happen in 100M chunks allocated to the **stress** program. Check both nodes as the container could be deployed to either. The next step will help if you have errors.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment hog
deployment.extensions/hog deleted

student@lfs458-node-1a0a:~$ kubectl apply -f hog.yaml
deployment.extensions/hog created
```

11. Should the resources not show as used, there may have been an issue inside of the container. Kubernetes shows it as running, but the actual workload has failed. Or the container may have failed; for example if you were missing a parameter the container may panic and show the following output.

```
student@lfs458-node-1a0a:~$ kubectl get pod
NAME                READY    STATUS    RESTARTS   AGE
hog-1985182137-5bz2w 0/1      Error     1           5s

student@lfs458-node-1a0a:~$ kubectl logs hog-1985182137-5bz2w
panic: cannot parse '150mi': unable to parse quantity's suffix

goroutine 1 [running]:
panic(0x5ff9a0, 0xc820014cb0)
    /usr/local/go/src/runtime/panic.go:481 +0x3e6
k8s.io/kubernetes/pkg/api/resource.MustParse(0x7ffe460c0e69, 0x5, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0)
    /usr/local/google/home/vishnuk/go/src/k8s.io/kubernetes/pkg/api/resource/quantity.go:134 +0x287
main.main()
    /usr/local/google/home/vishnuk/go/src/github.com/vishh/stress/main.go:24 +0x43
```

- Here is an example of an improper parameter. The container is running, but not allocating memory. It should show the usage requested from the YAML file.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY   STATUS    RESTARTS   AGE
hog-1603763060-x3vnn 1/1     Running   0           8s

student@lfs458-node-1a0a:~$ kubectl logs hog-1603763060-x3vnn
I0927 21:09:23.514921    1 main.go:26] Allocating "0" memory, in "4Ki" chunks, with a 1ms sleep \
                    between allocations
I0927 21:09:23.514984    1 main.go:39] Spawning a thread to consume CPU
I0927 21:09:23.514991    1 main.go:39] Spawning a thread to consume CPU
I0927 21:09:23.514997    1 main.go:29] Allocated "0" memory
```

Exercise 4.2: Resource Limits for a Namespace

The previous steps set limits for that particular deployment. You can also set limits on an entire namespace. We will create a new namespace and configure the **hog** deployment to run within. When set **hog** should not be able to use the previous amount of resources.

- Begin by creating a new namespace called `low-usage-limit` and verify it exists.

```
student@lfs458-node-1a0a:~$ kubectl create namespace low-usage-limit
namespace/low-usage-limit created

student@lfs458-node-1a0a:~$ kubectl get namespace
NAME                STATUS    AGE
default             Active    1h
kube-public         Active    1h
kube-system         Active    1h
low-usage-limit     Active    42s
```

- Create a YAML file which limits CPU and memory usage. The kind to use is `LimitRange`.

```
student@lfs458-node-1a0a:~$ vim low-resource-range.yaml

apiVersion: v1
kind: LimitRange
metadata:
  name: low-resource-range
spec:
  limits:
  - default:
      cpu: 1
      memory: 500Mi
    defaultRequest:
      cpu: 0.5
      memory: 100Mi
    type: Container
```

- Create the `LimitRange` object and assign it to the newly created namespace `low-usage-limit`

```
student@lfs458-node-1a0a:~$ kubectl create -f low-resource-range.yaml \
--namespace=low-usage-limit
limitrange/low-resource-range created
```

- Verify it works. Remember that every command needs a namespace and context to work. Defaults are used if not provided.

```
student@lfs458-node-1a0a:~$ kubectl get LimitRange
No resources found.
```

```
student@lfs458-node-1a0a:~$ kubectl get LimitRange --all-namespaces
NAMESPACE          NAME                CREATED AT
low-usage-limit     low-resource-range  2018-07-08T06:28:33Z
```

5. Create a new deployment in the namespace.

```
student@lfs458-node-1a0a:~$ kubectl create deployment limited-hog \
--image vish/stress -n low-usage-limit
deployment.apps/limited-hog created
```

6. List the current deployments. Note `hog` continues to run in the default namespace. If you chose to use the **Calico** network policy you may see a couple more than what is listed below.

```
student@lfs458-node-1a0a:~$ kubectl get deployments --all-namespaces
NAMESPACE          NAME           DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
default            hog            1         1         1             1           25m
kube-system        kube-dns       1         1         1             1           2d
low-usage-limit     limited-hog    1         1         1             1           1m
```

7. View all pods within the namespace. Remember you can use the **tab** key to complete the namespace. You may want to type the namespace first so that tab-completion is appropriate to that namespace instead of the default namespace.

```
student@lfs458-node-1a0a:~$ kubectl -n low-usage-limit get pods
NAME                                READY   STATUS    RESTARTS   AGE
limited-hog-2556092078-wnpnv        1/1     Running   0           3m
```

8. Look at the details of the pod. You will note it has the settings inherited from the entire namespace. The use of shell completion should work if you declare the namespace first.

```
student@lfs459-node-1a0a:~$ kubectl -n low-usage-limit get pod limited-hog-2556092078-wnpnv -o yaml
<output_omitted>
spec:
  containers:
  - image: vish/stress
    imagePullPolicy: Always
    name: stress
    resources:
      limits:
        cpu: "1"
        memory: 500Mi
      requests:
        cpu: 500m
        memory: 100Mi
    terminationMessagePath: /dev/termination-log
<output_omitted>
```

9. Copy and edit the config file for the original `hog` file. Add the namespace: line so that a new deployment would be in the `low-usage-limit` namespace.

```
student@lfs458-node-1a0a:~$ cp hog.yaml hog2.yaml

student@lfs458-node-1a0a:~$ vim hog2.yaml
....
  labels:
    app: hog
    name: hog
    namespace: low-usage-limit      #<<--- Add this line
    selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/hog
spec:
....
```

10. Open up extra terminal sessions so you can have **top** running in each. When the new deployment is created it will probably be scheduled on the node not yet under any stress.

Create the deployment.

```
student@lfs458-node-1a0a:~$ kubectl create -f hog2.yaml
deployment.extensions/hog created
```

11. View the deployments. Note there are two with the same name, but in different namespaces. You may also find the calico-typha deployment has no pods, nor has any requested. Our small cluster does not need to add **Calico** pods via this autoscaler.

```
student@lfs458-node-1a0a:~$ kubectl get deployments --all-namespaces
```

NAMESPACE	NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
default	hog	1	1	1	1	45m
kube-system	calico-typha	0	0	0	0	8h
kube-system	coredns	2	2	2	2	8h
low-usage-limit	hog	1	1	1	1	13s
low-usage-limit	limited-hog	1	1	1	1	5m

12. Look at the **top** output running in other terminals. You should find that both hog deployments are using about the same amount of resources, once the memory is fully allocated. Per-deployment settings override the global namespace settings. You should see something like the following lines one from each node, which indicates use of one processor and about 12 percent of your memory, were you on a system with 8G total.

```
25128 root      20   0 958532 954672   3180 R 100.0 11.7   0:52.27 stress
24875 root      20   0 958532 954800   3180 R 100.3 11.7  41:04.97 stress
```

13. Delete the hog deployments to recover system resources.

```
student@lfs458-node-1a0a:~$ kubectl -n low-usage-limit delete deployment hog
deployment.extensions "hog" deleted

student@lfs458-node-1a0a:~$ kubectl delete deployment hog
deployment.extensions "hog" deleted
```

Exercise 4.3: More Complex Deployment

We will now deploy a more complex demo application to test the cluster. When completed it will be a sock shopping site. The short URL is shown below for: <https://raw.githubusercontent.com/microservices-demo/microservices-demo/master/deploy/kubernetes/complete-demo.yaml>

1. Begin by downloading the pre-made YAML file from github.

```
student@lfs458-node-1a0a:~$ wget https://tinyurl.com/y8bn2awp -O complete-demo.yaml
Resolving tinyurl.com (tinyurl.com)... 104.20.218.42, 104.20.219.42,
Connecting to tinyurl.com (tinyurl.com)|104.20.218.42|:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://raw.githubusercontent.com/microservices-demo/microservices-dem...
--2017-11-02 16:54:27-- https://raw.githubusercontent.com/microservices-dem...
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.5...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101...
HTTP request sent, awaiting response... 200 OK
<output_omitted>
```

2. Find the expected namespaces inside the file. It should be `sock-shop`. Also note the various settings. This file will deploy several containers which work together, providing a shopping website. As we work with other parameters you could revisit this file to see potential settings.

```
student@lfs458-node-1a0a:~$ less complete-demo.yaml
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: carts-db
  labels:
    name: carts-db
  namespace: sock-shop
spec:
  replicas: 1
<output_omitted>
```

3. Create the namespace and verify it was made.

```
student@lfs458-node-1a0a:~$ kubectl create namespace sock-shop
namespace/sock-shop created
```

```
student@lfs458-node-1a0a:~$ kubectl get namespace
NAME                STATUS    AGE
default             Active    35m
kube-public         Active    35m
kube-system         Active    35m
low-usage-limit     Active    25m
sock-shop           Active    5s
```

4. View the images the new application will deploy.

```
student@lfs458-node-1a0a:~$ grep image complete-demo.yaml
  image: mongo
  image: weaveworksdemos/carts:0.4.8
  image: weaveworksdemos/catalogue-db:0.3.0
  image: weaveworksdemos/catalogue:0.3.5
  image: weaveworksdemos/front-end:0.3.12
  image: mongo
<output_omitted>
```

5. Create the new shopping website using the YAML file. Use the namespace you recently created. Note that the deployments match the images we saw in the file.

```
student@lfs458-node-1a0a:~$ kubectl apply -n sock-shop -f complete-demo.yaml
deployment "carts-db" created
service "carts-db" created
deployment "carts" created
service "carts" created
<output_omitted>
```

6. Using the proper namespace will be important. This can be set on a per-command basis or as a shell parameter. Note the first command shows no pods. We must remember to pass the proper namespace. Some containers may not have fully downloaded or deployed by the time you run the command.

```
student@lfs458-node-1a0a:~$ kubectl get pods
No resources found.
```

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop get pods
```

NAME	READY	STATUS	RESTARTS	AGE
carts-511261774-c4jwv	1/1	Running	0	71s
carts-db-549516398-tw9zs	1/1	Running	0	71s
catalogue-4293036822-sp5kt	1/1	Running	0	71s
catalogue-db-1846494424-qzhvk	1/1	Running	0	71s
front-end-2337481689-6s65c	1/1	Running	0	71s
orders-208161811-1gc6k	1/1	Running	0	71s
orders-db-2069777334-4sp01	1/1	Running	0	71s
payment-3050936124-2cn2l	1/1	Running	0	71s

```

queue-master-2067646375-vzq77    1/1    Running    0    71s
rabbitmq-241640118-vk3m9        0/1    ContainerCreating    0    71s
shipping-3132821717-lm7kn        0/1    ContainerCreating    0    71s
user-1574605338-24xrb            0/1    ContainerCreating    0    71s
user-db-2947298815-lx9kp         1/1    Running    0    71s

```

7. Verify the shopping cart is exposing a web page. Use the public IP address of your AWS node (not the one derived from the prompt) to view the page. Note the external IP is not yet configured. Find the NodePort service. First try port 80 then try port 30001 as shown under the PORTS column.

```

student@lfs458-node-1a0a:~$ kubectl get svc -n sock-shop
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
carts         ClusterIP     10.100.154.148 <none>         80/TCP           95s
carts-db      ClusterIP     10.111.120.73  <none>         27017/TCP        95s
catalogue     ClusterIP     10.100.8.203   <none>         80/TCP           95s
catalogue-db  ClusterIP     10.111.94.74   <none>         3306/TCP         95s
front-end     NodePort      10.98.2.137    <none>         80:30001/TCP     95s
orders        ClusterIP     10.110.7.215   <none>         80/TCP           95s
orders-db     ClusterIP     10.106.19.121  <none>         27017/TCP        95s
payment       ClusterIP     10.111.28.218  <none>         80/TCP           95s
queue-master  ClusterIP     10.102.181.253 <none>         80/TCP           95s
rabbitmq      ClusterIP     10.107.134.121 <none>         5672/TCP         95s
shipping      ClusterIP     10.99.99.127   <none>         80/TCP           95s
user          ClusterIP     10.105.126.10  <none>         80/TCP           95s
user-db       ClusterIP     10.99.123.228  <none>         27017/TCP        95s

```

8. Check to see which node is running the containers. Note that the webserver is answering on a node which is not hosting the all the containers. First we check the master, then the second node. The containers should have to do with **kube proxy** services and **calico**. The following is the **sudo docker ps** on both nodes.

```

student@lfs458-node-1a0a:~$ sudo docker ps
CONTAINER ID   IMAGE
d6b7353e5dc5   weaveworksdemos/user@sha256:2ffccc332963c89e035fea52201012208bf62df43a55fe461ad6598a5c757ab7
6c18f030f15b   weaveworksdemos/shipping@sha256:983305c948fde487f4a4acdeab5f898e89d577b4bc1ca3de7750076469ccad4
baaa8d67ebef   weaveworksdemos/queue-master@sha256:6292d3095f4c7aeed8d863527f8ef6d7a75d3128f20fc61e57f398c100142
<output_omitted>

```

```

student@lfs458-worker:~$ sudo docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED    STATUS    PORTS          NAMES
9452559caa0d   weaveworksdemos/payment@sha256:5ab1c9877480a018d4dda10d6dfa382776e6bca9fc1c60bacbb80903fde8cfe0
993017c7b476   weaveworksdemos/user-db@sha256:b43f0f8a76e0c908805fcec74d1ad7f4af4d93c4612632bd6dc20a87508e0b68
1356b0548ee8   weaveworksdemos/orders@sha256:b622e40e83433baf6374f15e076b53893f79958640fc6667dff597622eff03b9
<output_omitted>

```

9. Now we will shut down the shopping application. This can be done a few different ways. Begin by getting a listing of resources in all namespaces. There should be about 14 deployments.

```

student@lfs458-node-1a0a:~$ kubectl get deployment --all-namespaces
NAMESPACE     NAME          DESIRED    CURRENT    UP-TO-DATE    AVAILABLE    AGE
kube-system   calico-typha  0          0          0             0            33m
kube-system   coredns       2          2          2             2            33m
low-usage-limit limited-hog  1          1          1             1            33m
sock-shop     carts         1          1          1             1            6m44s
sock-shop     carts-db      1          1          1             1            6m44s
sock-shop     catalogue     1          1          1             1            6m44s
<output_omitted>

```

10. Use the terminal on the second node to get a count of the current docker containers. It should be something like 30, plus a line for status counted by **wc**. The main system should have something like 26 running, plus a line of status.

```

student@lfs458-node-1a0a:~$ sudo docker ps | wc -l
26

```



```
student@lfs458-worker:~$ sudo docker ps | wc -l
30
```

11. In order to complete maintenance we may need to move containers from a node and prevent new ones from deploying. One way to do this is to **drain**, or **cordon**, the node. Currently this will not affect DaemonSets, an object we will discuss in greater detail in the future. Begin by getting a list of nodes. Your node names will be different.

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE    VERSION
lfs458-worker       Ready    <none>   3h     v1.12.1
lfs458-node-1a0a    Ready    master   3h     v1.12.1
```

12. Modifying your second, worker node, update the node to **drain** the pods. Some resources may not drain, expect an error which we will work with next. Note the error includes aborting command which indicates the drain did not take place. Were you to check it would have the same number of containers running, but will show a new taint preventing the scheduler from assigning new pods.

```
student@lfs458-node-1a0a:~$ kubectl drain lfs458-worker
node/lfs458-worker cordoned
error: unable to drain node "lfs458-worker", aborting command...
```

```
There are pending nodes to be drained:
lfs458-worker
error: DaemonSet-managed pods (use --ignore-daemonsets to ignore):
calico-node-vndn7, kube-proxy-rjpls
```

```
student@lfs458-node-1a0a:~$ kubectl describe node |grep -i taint
Taints:                <none>
Taints:                node.kubernetes.io/unschedulable:NoSchedule
```

13. As the error output suggests we can use the **--ignore-daemonsets** options to ignore containers which are not intended to move. We will find a new error when we use this command, near the end of the output. The node will continue to have the same number of pods and containers running.

```
student@lfs458-node-1a0a:~$ kubectl drain lfs458-worker --ignore-daemonsets
node/worker cordoned
error: unable to drain node "lfs458-worker", aborting command...
```

```
There are pending nodes to be drained:
lfs458-worker
error: pods with local storage (use --delete-local-data to override):
carts-55f7f5c679-ffkq2, carts-db-5c55874946-w728d, orders-7b69bf5686-vtkcn
```

14. Run the command again. This time the output should both indicate the node has already been cordoned, then show the eviction of several pods. Not all pods will be gone as daemonsets will remain. Note the command is shown on two lines. You can omit the backslash and type the command on a single line.

```
student@lfs458-node-1a0a:~$ kubectl drain lfs458-worker \
--ignore-daemonsets --delete-local-data
```

```
node/lfs458-worker already cordoned
WARNING: Ignoring DaemonSet-managed pods: calico-node-vndn7, kube-proxy-rjpls; Deleting pods with local storage:
pod/carts-db-5c55874946-h42v2 evicted
pod/orders-db-7bc46bdb98-x5zrl evicted
pod/catalogue-db-66ff5bbbf5-2wmx4 evicted
pod/catalogue-5764fdf6d-8gk96 evicted
pod/orders-7b69bf5686-t82lz evicted
pod/front-end-f99dbcb9c-92q4p evicted
pod/carts-55f7f5c679-ppv7p evicted
```

15. Were you to look on your second, worker node, you would see there should be fewer pods and containers than before. These pods can only be evicted via a special taint which we will discuss in the scheduling chapter.

```
student@lfs458-worker:~$ sudo docker ps | wc -l
6
```

16. Update the node taint such that the scheduler will use the node again. Verify that no nodes have moved over to the worker node as the scheduler only checks when a pod is deployed.

```
student@lfs458-node-1a0a:~$ kubectl uncordon lfs458-worker
node/lfs458-worker uncordoned

student@lfs458-node-1a0a:~$ kubectl describe node |grep -i taint
Taints:          <none>
Taints:          <none>

student@lfs458-worker:~$ sudo docker ps | wc -l
6
```

17. As we clean up our sock shop let us see some differences between pods and deployments. Start with a list of the pods that are running in the sock-shop namespace.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop get pod
NAME                                READY    STATUS    RESTARTS   AGE
carts-db-549516398-tw9zs            1/1     Running   0          6h
catalogue-4293036822-sp5kt          1/1     Running   0          6h
<output_omitted>
```

18. Delete a few resources using the pod name.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop delete pod \
catalogue-4293036822-sp5kt catalogue-db-1846494424-qzhvk \
front-end-2337481689-6s65c orders-208161811-1gc6k \
orders-db-2069777334-4sp01
pod "catalogue-4293036822-sp5kt" deleted
pod "catalogue-db-1846494424-qzhvk" deleted
<output_omitted>
```

19. Check the status of the pods. There should be some pods running for only a few seconds. These will have the same name-stub as the Pods you recently deleted. The Deployment controller noticed expected number of Pods was not proper, so created new Pods until the current state matches the Pod manifest.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop get pod
NAME                                READY    STATUS    RESTARTS   AGE
catalogue-4293036822-mtz8m          1/1     Running   0          22s
catalogue-db-1846494424-16n2p        1/1     Running   0          22s
front-end-2337481689-6s65c          1/1     Terminating 0          6h
front-end-2337481689-80gwt          1/1     Running   0          22s
```

20. Delete some of the resources via deployments.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop delete deployment \
catalogue catalogue-db front-end orders
deployment "catalogue" deleted
deployment "catalogue-db" deleted
```

21. Check and both the pods and deployments you removed have not been recreated.

```
student@lfs458-node-1a0a:~$ kubectl -n sock-shop get pods |grep catalogue

student@lfs458-node-1a0a:~$ kubectl -n sock-shop get deployment
NAME            DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
carts           1         1         1             1          71m
carts-db        1         1         1             1          71m
orders-db       1         1         1             1          71m
payment         1         1         1             1          71m
```

queue-master	1	1	1	1	71m
rabbitmq	1	1	1	1	71m
shipping	1	1	1	1	71m
user	1	1	1	1	71m
user-db	1	1	1	1	71m

22. Delete the rest of the deployments. When no resources are found, examine the output of the `docker ps` command. None of the `sock-shop` containers should be found. Use the same file we created with to delete all of the objects made. You will get some errors because we deleted a few deployments by hand.

```
student@lfs458-node-1a0a:~$ kubectl delete -f complete-demo.yaml
<output_omitted>
```


Chapter 5

APIs and Access



5.1 Labs

Exercise 5.1: Configuring TLS Access

Overview

Using the Kubernetes API, **kubectl** makes API calls for you. With the appropriate TLS keys you could run **curl** as well use a **golang** client. Calls to the `kube-apiserver` get or set a `PodSpec`, or desired state. If the request represents a new state the **Kubernetes Control Plane** will update the cluster until the current state matches the specified state. Some end states may require multiple requests. For example, to delete a `ReplicaSet`, you would first set the number of replicas to zero, then delete the `ReplicaSet`.

An API request must pass information as JSON. **kubectl** converts `.yaml` to JSON when making an API request on your behalf. The API request has many settings, but must include `apiVersion`, `kind` and `metadata`, and `spec` settings to declare what kind of container to deploy. The `spec` fields depend on the object being created.

We will begin by configuring remote access to the `kube-apiserver` then explore more of the API.

Configuring TLS Access

1. Begin by reviewing the **kubectl** configuration file. We will use the three certificates and the API server address.

```
student@lfs458-node-1a0a:~$ less ~/.kube/config
<output_omitted>
```

2. We will set the certificates as variables. You may want to double-check each parameter as you set it. Begin with setting the `client-certificate-data` key.

```
student@lfs458-node-1a0a:~$ export client=$(grep client-cert ~/.kube/config |cut -d" " -f 6)

student@lfs458-node-1a0a:~$ echo $client
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUM4akNDQWRxZ0F3SUJ
```

```
BZ01JRy9wbC9rWEpNdm3RFFZSktvWklodmNOQVFFTEJRQXdGVEVUTUJFR0
ExVUUKQXhNS2EzVmlaWEp1WlhSbGN6QWVGdzB4TnpFeU1UTXh0elEyTXpKY
UZ3MHhPREV5TVRNeE56UTJNelJhTURReApGekFWQmdOVkJBb1REbk41YzNS
<output_omitted>
```

3. Almost the same command, but this time collect the `client-key-data` as the `key` variable.

```
student@lfs458-node-1a0a:~$ export key=$(grep client-key-data ~/.kube/config |cut -d " " -f 6)

student@lfs458-node-1a0a:~$ echo $key
<output_omitted>
```

4. Finally set the `auth` variable with the `certificate-authority-data` key.

```
student@lfs458-node-1a0a:~$ export auth=$(grep certificate-authority-data ~/.kube/config |cut -d " " -f 6)

student@lfs458-node-1a0a:~$ echo $auth
<output_omitted>
```

5. Now encode the keys for use with **curl**.

```
student@lfs458-node-1a0a:~$ echo $client | base64 -d - > ./client.pem

student@lfs458-node-1a0a:~$ echo $key | base64 -d - > ./client-key.pem

student@lfs458-node-1a0a:~$ echo $auth | base64 -d - > ./ca.pem
```

6. Pull the API server URL from the config file.

```
student@lfs458-node-1a0a:~$ kubectl config view |grep server
server: https://10.128.0.3:6443
```

7. Use **curl** command and the encoded keys to connect to the API server.

```
student@lfs458-node-1a0a:~$ curl --cert ./client.pem \
--key ./client-key.pem \
--cacert ./ca.pem \
https://10.128.0.3:6443/api/v1/pods

{
  "kind": "PodList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/pods",
    "resourceVersion": "239414"
  },
  <output_omitted>
}
```

8. If the previous command was successful, create a JSON file to create a new pod. Remember to look for this file in the tarball output, it can save you some typing.

```
student@lfs458-node-1a0a:~$ vim curlpod.json
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    "namespace": "default",
    "labels": {
      "name": "examplepod"
    }
  },
  "spec": {
    "containers": [{
```

```

        "name": "nginx",
        "image": "nginx",
        "ports": [{"containerPort": 80}]
    }
}

```

- The previous **curl** command can be used to build a XPOST API call. There will be a lot of output, including the scheduler and taints involved. Read through the output. In the last few lines the phase will probably show Pending, as it's near the beginning of the creation process.

```

student@lfs458-node-1a0a:~$ curl --cert ./client.pem \
--key ./client-key.pem --cacert ./ca.pem \
https://10.128.0.3:6443/api/v1/namespaces/default/pods \
-XPOST -H'Content-Type: application/json' \
-d@curlpod.json

{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    <output_omitted>
  }
}

```

- Verify the new pod exists and shows a Running status.

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
curlpod       1/1     Running   0           45s

```

Exercise 5.2: Explore API Calls

- One way to view what a command does on your behalf is to use **strace**. In this case, we will look for the current endpoints, or targets of our API calls.

```

student@lfs458-node-1a0a:~$ kubectl get endpoints
NAME          ENDPOINTS          AGE
kubernetes    10.128.0.3:6443    3h

```

- Run this command again, preceded by **strace**. You will get a lot of output. Near the end you will note several **openat** functions to a local directory, `/home/student/.kube/cache/discovery/10.128.0.3_6443`. If you cannot find the lines, you may want to redirect all output to a file and **grep** for them.

```

student@lfs458-node-1a0a:~$ strace kubectl get endpoints
execve("/usr/bin/kubectl", ["kubectl", "get", "endpoints"], [/*...
....
openat(AT_FDCWD, "/home/student/.kube/cache/discovery/10.128.0.3_6443..
<output_omitted>

```

- Change to the parent directory and explore. Your endpoint IP will be different, so replace the following with one suited to your system.

```

student@lfs458-node-1a0a:~$ cd /home/student/.kube/cache/discovery/

student@lfs458-node-1a0a:~/.kube/cache/discovery$ ls
10.128.0.3_6443

student@lfs458-node-1a0a:~/.kube/cache/discovery$ cd 10.128.0.3_6443/

```

- View the contents. You will find there are directories with various configuration information for kubernetes.

```
student@lfs458-node-1a0a: ~/.kube/cache/discovery/10.128.0.3_6443$ ls
admissionregistration.k8s.io  batch          policy
apiextensions.k8s.io         certificates.k8s.io  rbac.authorization.k8s.io
apiregistration.k8s.io      coordination.k8s.io  scheduling.k8s.io
apps                        crd.projectcalico.org  servergroups.json
authentication.k8s.io       events.k8s.io      storage.k8s.io
authorization.k8s.io        extensions        v1
autoscaling                networking.k8s.io
```

5. Use the find command to list out the subfiles. The prompt has been modified to look better on this page.

```
student@lfs458-node-1a0a: ./10.128.0.3_6443$ find .
.
./events.k8s.io
./events.k8s.io/v1beta1
./events.k8s.io/v1beta1/serverresources.json
./apps
./apps/v1
./apps/v1/serverresources.json
./apps/v1beta1
./apps/v1beta1/serverresources.json
<output_omitted>
```

6. View the objects available in version 1 of the API. For each object, or kind:, you can view the verbs or actions for that object, such as create seen in the following example. Note the prompt has been truncated for the command to fit on one line. Some are HTTP verbs, such as GET, others are product specific options, not standard HTTP verbs.

```
student@lfs458-node-1a0a:.$ python -m json.tool v1/serverresources.json
{
  "apiVersion": "v1",
  "groupVersion": "v1",
  "kind": "APIResourceList",
  "resources": [
    {
      "kind": "Binding",
      "name": "bindings",
      "namespaced": true,
      "singularName": "",
      "verbs": [
        "create"
      ]
    },
    ...
  ],
  ...
}
<output_omitted>
```

7. Some of the objects have shortNames, which makes using them on the command line much easier. Locate the shortName for endpoints.

```
student@lfs458-node-1a0a:.$ python -m json.tool v1/serverresources.json | less
.
{
  "kind": "Endpoints",
  "name": "endpoints",
  "namespaced": true,
  "shortNames": [
    "ep"
  ],
  "singularName": "",
  "verbs": [
    "create",
    "delete",
    ...
  ]
}
.
```

8. Use the shortName to view the endpoints. It should match the output from the previous command.


```
student@lfs458-node-1a0a:.$ kubectl get ep
NAME           ENDPOINTS          AGE
kubernetes     10.128.0.3:6443    3h
```

9. We can see there are 37 objects in version 1 file.

```
student@lfs458-node-1a0a:.$ python -m json.tool v1/serverresources.json | grep kind
"kind": "APIResourceList",
    "kind": "Binding",
    "kind": "ComponentStatus",
    "kind": "ConfigMap",
    "kind": "Endpoints",
    "kind": "Event",
<output_omitted>
```

10. Looking at another file we find nine more.

```
student@lfs458-node-1a0a:.$ python -m json.tool \
    apps/v1beta1/serverresources.json | grep kind
"kind": "APIResourceList",
    "kind": "ControllerRevision",
    "kind": "Deployment",
<output_omitted>
```

11. Delete the curlpod to recoup system resources.

```
student@lfs458-node-1a0a:.$ kubectl delete po curlpod
pod "curlpod" deleted
```

12. Take a look around the other files in this directory as time permits.

Chapter 6

API Objects



6.1 Labs

Exercise 6.1: RESTful API Access

Overview

We will continue to explore ways of accessing the control plane of our cluster. In the security chapter we will discuss there are several authentication methods, one of which is use of a Bearer token. We will work with one then deploy a local proxy server for application-level access to the Kubernetes API.

RESTful API Access

We will use the **curl** command to make API requests to the cluster, in an in-secure manner. Once we know the IP address and port, then the token we can retrieve cluster data in a RESTful manner. By default most of the information is restricted, but changes to authentication policy could allow more access.

1. First we need to know the IP and port of a node running a replica of the API server. The master system will typically have one running. Use **kubectl config view** to get overall cluster configuration, and find the server entry. This will give us both the IP and the port.

```
student@lfs458-node-1a0a:~$ kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
    server: https://10.128.0.3:6443
  name: kubernetes
<output_omitted>
```

2. Next we need to find the bearer token. This is part of a default token. Look at a list of tokens, first all on the cluster, then just those in the default namespace. There will be a secret for each of the controllers of the cluster.

- Pods can also make use of included certificates to use the API. The certificates are automatically made available to a pod under the `/var/run/secrets/kubernetes.io/serviceaccount/`. We will deploy a simple Pod and view the resources. If you view the `token` file you will find it is the same value we put into the `$token` variable. The `-i` will request a `-t` terminal session of the busybox container. Once you exit the container will not restart and the pod will show as completed.

```
student@lfs458-node-1a0a:~$ kubectl run -i -t busybox --image=busybox \
--restart=Never
# ls /var/run/secrets/kubernetes.io/serviceaccount/
ca.crt      namespace  token
/ # exit
```

Exercise 6.2: Using the Proxy

Another way to interact with the API is via a proxy. The proxy can be run from a node or from within a Pod through the use of a sidecar. In the following steps we will deploy a proxy listening to the loopback address. We will use `curl` to access the API server. If the `curl` request works, but does not from outside the cluster, we have narrowed down the issue to authentication and authorization instead of issues further along the API ingestion process.

- Begin by starting the proxy. It will start in the foreground by default. There are several options you could pass. Begin by reviewing the help output.

```
student@lfs458-node-1a0a:~$ kubectl proxy -h
Creates a proxy server or application-level gateway between localhost
and the Kubernetes API Server. It also allows serving static content
over specified HTTP path. All incoming data enters through one port
and gets forwarded to the remote kubernetes API Server port, except
for the path matching the static content path.
```

Examples:

```
# To proxy all of the kubernetes api and nothing else, use:
```

```
$ kubectl proxy --api-prefix=/
<output_omitted>
```

- Start the proxy while setting the API prefix, and put it in the background. You may need to use `enter` to view the prompt.

```
student@lfs458-node-1a0a:~$ kubectl proxy --api-prefix=/ &
[1] 22500
Starting to serve on 127.0.0.1:8001
```

- Now use the same `curl` command, but point toward the IP and port shown by the proxy. The output should be the same as without the proxy, but may be formatted differently.

```
student@lfs458-node-1a0a:~$ curl http://127.0.0.1:8001/api/
<output_omitted>
```

- Make an API call to retrieve the namespaces. The command did not work in the previous section due to permissions, but should work now as the proxy is making the request on your behalf.

```
student@lfs458-node-1a0a:~$ curl http://127.0.0.1:8001/api/v1/namespaces
{
  "kind": "NamespaceList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/namespaces",
    "resourceVersion": "86902"
  }
}
<output_omitted>
```

Exercise 6.3: Working with Jobs

While most API objects are deployed such that they continue to be available there are some which we may want to run a particular number of times called a Job, and others on a regular basis called a CronJob

Create A Job

1. Create a job which will run a container which sleeps for three seconds then stops.

```
student@lfs458-node-1a0a:~$ vim job.yaml
apiVersion: batch/v1
kind: Job
metadata:
  name: sleepy
spec:
  template:
    spec:
      containers:
      - name: resting
        image: busybox
        command: ["/bin/sleep"]
        args: ["3"]
        restartPolicy: Never
```

2. Create the job, then verify and view the details. The example shows checking the job three seconds in and then again after it has completed. You may see different output depending on how fast you type.

```
student@lfs458-node-1a0a:~$ kubectl create -f job.yaml
job.batch/sleepy created

student@lfs458-node-1a0a:~$ kubectl get job
NAME          COMPLETIONS  DURATION  AGE
sleepy        0/1          3s        3s

student@lfs458-node-1a0a:~$ kubectl describe jobs.batch sleepy
Name:          sleepy
Namespace:     default
Selector:      controller-uid=24c91245-d0fb-11e8-947a-42010a800002
Labels:        controller-uid=24c91245-d0fb-11e8-947a-42010a800002
               job-name=sleepy
Annotations:    <none>
Parallelism:   1
Completions:   1
Start Time:    Tue, 16 Oct 2018 04:22:50 +0000
Completed At:  Tue, 16 Oct 2018 04:22:55 +0000
Duration:      5s
Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl get job
NAME          COMPLETIONS  DURATION  AGE
sleepy        1/1          5s        17s
```

3. View the configuration information of the job. There are three parameters we can use to affect how the job runs. Use **-o yaml** to see these parameters. We can see that backoffLimit, completions, and the parallelism. We'll add these parameters next.

```
student@lfs458-node-1a0a:~$ kubectl get jobs.batch sleepy -o yaml

<output_omitted>
uid: c2c3a80d-d0fc-11e8-947a-42010a800002
spec:
  backoffLimit: 6
  completions: 1
  parallelism: 1
  selector:
    matchLabels:
<output_omitted>
```

4. As the job continues to AGE in a completion state, delete the job.

```
student@lfs458-node-1a0a:~$ kubectl delete jobs.batch sleepy
job.batch "sleepy" deleted
```

5. Edit the YAML and add the `completions:` parameter and set it to 5.

```
student@lfs458-node-1a0a:~$ vim job.yaml
<output_omitted>
metadata:
  name: sleepy
spec:
  completions: 5    #<--Add this line
  template:
    spec:
      containers:
<output_omitted>
```

6. Create the job again. As you view the job note that `COMPLETIONS` begins as zero of 5.

```
student@lfs458-node-1a0a:~$ kubectl create -f job.yaml
job.batch/sleepy created
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs.batch
NAME          COMPLETIONS  DURATION  AGE
sleepy        0/5           5s        5s
```

7. View the pods that running. Again the output may be different depending on the speed of typing.

```
student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY  STATUS      RESTARTS  AGE
sleepy-z5tnh   0/1    Completed    0          8s
sleepy-zd692   1/1    Running      0          3s
<output_omitted>
```

8. Eventually all the jobs will have completed. Verify then delete the job.

```
student@lfs458-node-1a0a:~$ kubectl get jobs
NAME          COMPLETIONS  DURATION  AGE
sleepy        5/5           26s       10m

student@lfs458-node-1a0a:~$ kubectl delete jobs.batch sleepy
job.batch "sleepy" deleted
```

9. Edit the YAML again. This time add in the `parallelism:` parameter. Set it to 2 such that two pods at a time will be deployed.

```
student@lfs458-node-1a0a:~$ vim job.yaml
<output_omitted>
  name: sleepy
spec:
  completions: 5
  parallelism: 2    #<-- Add this line
  template:
    spec:
<output_omitted>
```

10. Create the job again. You should see the pods deployed two at a time until all five have completed.

```
student@lfs458-node-1a0a:~$ kubectl create -f job.yaml
job.batch/sleepy created
```

```
student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY  STATUS      RESTARTS  AGE
```

```
sleepy-8xwpc          1/1      Running   0          5s
sleepy-xjqnf          1/1      Running   0          5s
<output_omitted>
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs
NAME      COMPLETIONS  DURATION  AGE
sleepy    3/5           11s       11s
```

11. Add a parameter which will stop the job after a certain number of seconds. Set the `activeDeadlineSeconds:` to 15. The job and all pods will end once it runs for 15 seconds. We will also increase the sleep argument to five, just to be sure does not expire by itself.

```
student@lfs458-node-1a0a:~$ vim job.yaml
<output_omitted>
  completions: 5
  parallelism: 2
  activeDeadlineSeconds: 15  #<-- Add this line
  template:
    spec:
      containers:
      - name: resting
        image: busybox
        command: ["/bin/sleep"]
        args: ["5"]          #<-- Edit this line
<output_omitted>
```

12. Delete and recreate the job again. It should run for 15 seconds, usually 3/5, then continue to age without further completions.

```
student@lfs458-node-1a0a:~$ kubectl delete jobs.batch sleepy
job.batch "sleepy" deleted

student@lfs458-node-1a0a:~$ kubectl create -f job.yaml
job.batch/sleepy created

student@lfs458-node-1a0a:~$ kubectl get jobs
NAME      COMPLETIONS  DURATION  AGE
sleepy    1/5           6s        6s

student@lfs458-node-1a0a:~$ kubectl get jobs
NAME      COMPLETIONS  DURATION  AGE
sleepy    3/5           16s       16s
```

13. View the `message:` entry in the `Status` section of the object YAML output.

```
student@lfs458-node-1a0a:~$ kubectl get job sleepy -o yaml
<output_omitted>
status:
  conditions:
  - lastProbeTime: 2018-10-16T05:45:14Z
    lastTransitionTime: 2018-10-16T05:45:14Z
    message: Job was active longer than specified deadline
    reason: DeadlineExceeded
    status: "True"
    type: Failed
  failed: 2
  startTime: 2018-10-16T05:44:59Z
  succeeded: 3
```

14. Delete the job.

```
student@lfs458-node-1a0a:~$ kubectl delete jobs.batch sleepy
job.batch "sleepy" deleted
```


Create a CronJob

A CronJob creates a watch loop which will create a batch job on your behalf when the time becomes true. We Will use our existing Job file to start.

1. Copy the Job file to a new file.

```
student@lfs458-node-1a0a:~$ cp job.yaml cronjob.yaml
```

2. Edit the file to look like the annotated file shown below. Edit the lines mentioned below. The three parameters we added will need to be removed. Other lines will need to be further indented.

```
student@lfs458-node-1a0a:~$ vim cronjob.yaml
apiVersion: batch/v1beta1    #<-- Add beta1 to be v1beta1
kind: CronJob                #<-- Update this line to CronJob
metadata:
  name: sleepy
spec:
  schedule: "*/2 * * * *"    #<-- Add Linux style cronjob syntax
  jobTemplate:               #<-- New jobTemplate and spec move
    spec:
      template:              #<-- This and following lines move
        spec:                #<-- four spaces to the right
          containers:
            - name: resting
              image: busybox
              command: ["/bin/sleep"]
              args: ["3"]
          restartPolicy: Never
```

3. Create the new CronJob. View the jobs. It will take two minutes for the CronJob to run and generate a new batch Job.

```
student@lfs458-node-1a0a:~$ kubectl create -f cronjob.yaml
cronjob.batch/sleepy created
```

```
student@lfs458-node-1a0a:~$ kubectl get cronjobs.batch
NAME      SCHEDULE    SUSPEND   ACTIVE   LAST SCHEDULE   AGE
sleepy    */2 * * * *   False     0        <none>          8s
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs.batch
No resources found.
```

4. After two minutes you should see jobs start to run.

```
student@lfs458-node-1a0a:~$ kubectl get cronjobs.batch
NAME      SCHEDULE    SUSPEND   ACTIVE   LAST SCHEDULE   AGE
sleepy    */2 * * * *   False     0        21s             2m1s
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs.batch
NAME                COMPLETIONS   DURATION   AGE
sleepy-1539722040    1/1           5s         18s
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs.batch
NAME                COMPLETIONS   DURATION   AGE
sleepy-1539722040    1/1           5s         5m17s
sleepy-1539722160    1/1           6s         3m17s
sleepy-1539722280    1/1           6s         77s
```

5. Ensure that if the job continues for more than 10 seconds it is terminated. We will first edit the **sleep** command to run for 30 seconds then add the `activeDeadlineSeconds`: entry to the container.

```
student@lfs458-node-1a0a:~$ vim cronjob.yaml
....
jobTemplate:
  spec:
    template:
      spec:
        activeDeadlineSeconds: 10 #<-- Add this line
        containers:
        - name: resting
....
```

6. Delete and recreate the CronJob. It may take a couple of minutes for the batch Job to be created and terminate due to the timer.

```
student@lfs458-node-1a0a:~$ kubectl delete cronjobs.batch sleepy
cronjob.batch "sleepy" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl create -f cronjob.yaml
cronjob.batch/sleepy created
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs
NAME                COMPLETIONS  DURATION  AGE
sleepy-1539723240   0/1           61s       61s
```

```
student@lfs458-node-1a0a:~$ kubectl get cronjobs.batch
NAME    SCHEDULE    SUSPEND  ACTIVE  LAST SCHEDULE  AGE
sleepy  */2 * * * *  False    1       72s           94s
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs
NAME                COMPLETIONS  DURATION  AGE
sleepy-1539723240   0/1           75s       75s
```

```
student@lfs458-node-1a0a:~$ kubectl get jobs
NAME                COMPLETIONS  DURATION  AGE
sleepy-1539723240   0/1           2m19s     2m19s
sleepy-1539723360   0/1           19s       19s
```

```
student@lfs458-node-1a0a:~$ kubectl get cronjobs.batch
NAME    SCHEDULE    SUSPEND  ACTIVE  LAST SCHEDULE  AGE
sleepy  */2 * * * *  False    2       31s           2m53s
```

7. Clean up by deleting the CronJob.

```
student@lfs458-node-1a0a:~$ kubectl delete cronjobs.batch sleepy
cronjob.batch "sleepy" deleted
```

Chapter 7

Managing State With Deployments



7.1 Labs

Exercise 7.1: Working with ReplicaSets

Overview

Understanding and managing the state of containers is a core Kubernetes task. In this lab we will first explore the API objects used to manage groups of containers. The objects available have changed as Kubernetes has matured, so the Kubernetes version in use will determine which are available. Our first object will be a `ReplicaSet`, which does not include newer management features found with `Deployments`. A `Deployment` will manage `ReplicaSets` for you. We will also work with another object called a `DaemonSet` which ensures a container is running on newly added node.

Then we will update the software in a container, view the revision history, and roll-back to a previous version.

Working with ReplicaSets

A `ReplicaSet` is a next-generation of a `Replication Controller`, which differs only in the selectors supported. The only reason to use a `ReplicaSet` anymore is if you have no need for updating container software or require update orchestration which won't work with the typical process.

1. View any current `ReplicaSets`. If you deleted resources at the end of a previous lab, you should have none reported in the default namespace.

```
student@lfs458-node-1a0a:~$ kubectl get rs
No resources found.
```

2. Create a YAML file for a simple `ReplicaSet`. The `apiVersion` setting depends on the version of Kubernetes you are using. Versions 1.8 and beyond will use `apps/v1beta1`, then perhaps someday `apps/v1beta2` and then probably a stable `apps/v1`. We will use an older version of **nginx** then update to a newer version later in the exercise.

```
student@lfs458-node-1a0a:~$ vim rs.yaml
```

```

apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: rs-one
spec:
  replicas: 2
  template:
    metadata:
      labels:
        system: ReplicaOne
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80

```

3. Create the ReplicaSet:

```

student@lfs458-node-1a0a:~$ kubectl create -f rs.yaml
replicaset.extensions/rs-one created

```

4. View the newly created ReplicaSet:

```

student@lfs458-node-1a0a:~$ kubectl describe rs rs-one
Name:          rs-one
Namespace:     default
Selector:      system=ReplicaOne
Labels:        system=ReplicaOne
Annotations:   <none>
Replicas:      2 current / 2 desired
Pods Status:   2 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:       system=ReplicaOne
  Containers:
    nginx:
      Image:      nginx:1.7.9
      Port:       80/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:      <none>
  Events:        <none>

```

5. View the Pods created with the ReplicaSet. From the yaml file created there should be two Pods. You may see a Completed busybox which will be cleared out eventually.

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
rs-one-2p9x4   1/1     Running   0           5m4s
rs-one-3c6pb   1/1     Running   0           5m4s

```

6. Now we will delete the ReplicaSet, but not the Pods it controls.

```

student@lfs458-node-1a0a:~$ kubectl delete rs rs-one --cascade=false
replicaset.extensions "rs-one" deleted

```

View the ReplicaSet and Pods again:

7. student@lfs458-node-1a0a:~\$ kubectl describe rs rs-one

Error from server (NotFound): replicaset.extensions "rs-one" not found

```

student@lfs458-node-1a0a:~$ kubectl get pods

```

NAME	READY	STATUS	RESTARTS	AGE
rs-one-2p9x4	1/1	Running	0	7m
rs-one-3c6pb	1/1	Running	0	7m

8. Create the ReplicaSet again. As long as we do not change the selector field, the new ReplicaSet should take ownership. Pod software versions cannot be updated this way.

```
student@lfs458-node-1a0a:~$ kubectl create -f rs.yaml
replicaset.extensions/rs-one created
```

9. View the age of the ReplicaSet and then the Pods within:

```
student@lfs458-node-1a0a:~$ kubectl get rs
NAME          DESIRED  CURRENT  READY  AGE
rs-one        2        2        2      46s

student@lfs458-node-1a0a:~$ kubectl get pods
NAME          READY  STATUS   RESTARTS  AGE
rs-one-2p9x4  1/1    Running  0         8m
rs-one-3c6pb  1/1    Running  0         8m
```

10. We will now isolate a Pod from its ReplicaSet. Begin by editing the label of a Pod. We will change the system: parameter to be IsolatedPod.

```
student@lfs458-node-1a0a:~$ kubectl edit po rs-one-3c6pb

....
  labels:
    system: IsolatedPod  #<-- Change from ReplicaOne
    name: rs-one-3c6pb
....
```

11. View the number of pods within the ReplicaSet. You should see two running.

```
student@lfs458-node-1a0a:~$ kubectl get rs
NAME          DESIRED  CURRENT  READY  AGE
rs-one        2        2        2      4m
```

12. Now view the pods with the label key of system. You should note that there are three, with one being newer than others. The ReplicaSet made sure to keep two replicas, replacing the Pod which was isolated.

```
student@lfs458-node-1a0a:~$ kubectl get po -L system
NAME          READY  STATUS   RESTARTS  AGE    SYSTEM
rs-one-3c6pb  1/1    Running  0         10m    IsolatedPod
rs-one-2p9x4  1/1    Running  0         10m    ReplicaOne
rs-one-dq5xd  1/1    Running  0         30s    ReplicaOne
```

13. Delete the ReplicaSet, then view any remaining Pods.

```
student@lfs458-node-1a0a:~$ kubectl delete rs rs-one
replicaset.extensions "rs-one" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY  STATUS   RESTARTS  AGE
rs-one-3c6pb  1/1    Running  0         14m
rs-one-dq5xd  0/1    Terminating  0         4m
```

14. In the above example the Pods had not finished termination. Wait for a bit and check again. There should be no ReplicaSets, but one Pod.

```
student@lfs458-node-1a0a:~$ kubectl get rs
No resources found.

student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY   STATUS    RESTARTS   AGE
rs-one-3c6pb  1/1     Running   0           16m
```

15. Delete the remaining Pod using the label.

```
student@lfs458-node-1a0a:~$ kubectl delete po -l system=IsolatedPod
pod "rs-one-3c6pb" deleted
```

Exercise 7.2: Working with DaemonSets

A DaemonSet is a watch loop object like a Deployment which we have been working with in the rest of the labs. The DaemonSet ensures that when a node is added to a cluster a pods will be created on that node. A Deployment would only ensure a particular number of pods are created in general, several could be on a single node. Using a DaemonSet can be helpful to ensure applications are on each node, helpful for things like metrics and logging especially in large clusters where hardware may be swapped out often. Should a node be removed from a cluster the DaemonSet would ensure the Pods are garbage collected before removal. Starting with Kubernetes v1.12 the scheduler handles DaemonSet deployment which means we can now configure certain nodes to not have a particular DaemonSet pods.

This extra step of automation can be useful for using with products like **ceph** where storage is often added or removed, but perhaps among a subset of hardware. They allow for complex deployments when used with declared resources like memory, CPU or volumes.

1. We begin by creating a yaml file. In this case the kind would be set to DaemonSet. For ease of use we will copy the previously created `rs.yaml` file and make a couple edits. Remove the Replicas: 2 line.

```
student@lfs458-node-1a0a:~$ cp rs.yaml ds.yaml

student@lfs458-node-1a0a:~$ vim ds.yaml
....
kind: DaemonSet
....
  name: ds-one
....
  replicas: 2 #<<<----Remove this line
....
  system: DaemonSetOne
....
```

2. Create and verify the newly formed DaemonSet. There should be one Pod per node in the cluster.

```
student@lfs458-node-1a0a:~$ kubectl create -f ds.yaml
daemonset.extensions/ds-one created

student@lfs458-node-1a0a:~$ kubectl get ds
NAME    DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE-SELECTOR   AGE
ds-one  2         2         2       2            2          <none>          1m

student@lfs458-node-1a0a:~$ kubectl get po
NAME          READY   STATUS    RESTARTS   AGE
ds-one-b1dcv  1/1     Running   0           2m
ds-one-z31r4  1/1     Running   0           2m
```

3. Verify the image running inside the Pods. We will use this information in the next section.

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-one-b1dcv | grep Image:
Image:          nginx:1.7.9
```

Exercise 7.3: Rolling Updates and Rollbacks

One of the advantages of micro-services is the ability to replace and upgrade a container while continuing to respond to client requests. We will use the default `OnDelete` setting that upgrades a container when the predecessor is deleted, then the use the `RollingUpdate` feature as well.

1. Begin by viewing the current `updateStrategy` setting for the `DaemonSet` created in the previous section.

```
student@lfs458-node-1a0a:~$ kubectl get ds ds-one -o yaml \
  | grep -A 1 Strategy
updateStrategy:
  type: OnDelete
```

2. Update the `DaemonSet` to use a newer version of the **nginx** server. This time use the **set** command instead of **edit**. Set the version to be `1.8.1-alpine`.

```
student@lfs458-node-1a0a:~$ kubectl set image ds ds-one nginx=nginx:1.8.1-alpine
daemonset.extensions/ds-one image updated
```

3. Verify that the `Image:` parameter for the Pod checked in the previous section is unchanged.

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-one-b1dcv |grep Image:
Image:                nginx:1.7.9
```

4. Delete the Pod. Wait until the replacement Pod is running and check the version.

```
student@lfs458-node-1a0a:~$ kubectl delete po ds-one-b1dcv
pod "ds-one-b1dcv" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY   STATUS    RESTARTS   AGE
ds-one-xc86w        1/1     Running   0           19s
ds-one-z31r4        1/1     Running   0           4m8s
```

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-one-xc86w |grep Image:
Image:                nginx:1.8.1-alpine
```

5. View the image running on the older Pod. It should still show version 1.7.9.

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-one-z31r4 |grep Image:
Image:                nginx:1.7.9
```

6. View the history of changes for the `DaemonSet`. You should see two revisions listed. The number of revisions kept is set in the `DaemonSet` with `v.1.12.1` the history kept has increased to ten from two, by default.

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds ds-one
daemonsets "ds-one"
REVISION    CHANGE-CAUSE
1           <none>
2           <none>
```

7. View the settings for the various versions of the `DaemonSet`. The `Image:` line should be the only difference between the two outputs.

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds ds-one --revision=1
daemonsets "ds-one" with revision #1
Pod Template:
  Labels:            system=DaemonSetOne
  Containers:
    nginx:
      Image:          nginx:1.7.9
      Port:           80/TCP
      Environment:    <none>
```

```
Mounts:      <none>
Volumes:     <none>
```

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds ds-one --revision=2
....
Image:      nginx:1.8.1-alpine
.....
```

8. Use `kubectl rollout undo` to change the DaemonSet back to an earlier version. As we are still using the `OnDelete` strategy there should be no change to the Pods.

```
student@lfs458-node-1a0a:~$ kubectl rollout undo ds ds-one --to-revision=1
daemonset.extensions/ds-one rolled back
```

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-one-xc86w |grep Image:
Image:      nginx:1.8.1-alpine
```

9. Delete the Pod, wait for the replacement to spawn then check the image version again.

```
student@lfs458-node-1a0a:~$ kubectl delete po ds-one-xc86w
pod "ds-one-xc86w" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY   STATUS    RESTARTS   AGE
ds-one-qc72k        1/1     Running   0           10s
ds-one-xc86w        0/1     Terminating 0           12m
ds-one-z31r4        1/1     Running   0           28m
```

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-one-qc72k |grep Image:
Image:      nginx:1.7.9
```

10. View the details of the DaemonSet. The Image should be v1.7.9 in the output.

```
student@lfs458-node-1a0a:~$ kubectl describe ds |grep Image:
Image:      nginx:1.7.9
```

11. View the current configuration for the DaemonSet in YAML output. Look for the update strategy near the end of the output.

```
student@lfs458-node-1a0a:~$ kubectl get ds ds-one -o yaml
apiVersion: extensions/v1beta1
kind: DaemonSet
....
  terminationGracePeriodSeconds: 30
  templateGeneration: 3
  updateStrategy:
    type: OnDelete
status:
  currentNumberScheduled: 2
....
```

12. Create a new DaemonSet, this time setting the update policy to `RollingUpdate`. Begin by generating a new config file.

```
student@lfs458-node-1a0a:~$ kubectl get ds ds-one -o yaml --export > ds2.yaml
```

13. Edit the file. Change the name, around line eight and the update strategy around line 38.

```
student@lfs458-node-1a0a:~$ vim ds2.yaml
....
  name: ds-two
....
  type: RollingUpdate
```


14. Create the new DaemonSet and verify the **nginx** version in the new pods.

```
student@lfs458-node-1a0a:~$ kubectl create -f ds2.yaml
daemonset.extensions/ds-two created

student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY    STATUS    RESTARTS   AGE
ds-one-qc72k        1/1      Running   0           28m
ds-one-z31r4        1/1      Running   0           57m
ds-two-10khc        1/1      Running   0           5m
ds-two-kzp9g        1/1      Running   0           5m

student@lfs458-node-1a0a:~$ kubectl describe po ds-two-10khc |grep Image:
Image:                nginx:1.7.9
```

15. Edit the configuration file and set the image to a newer version such as 1.8.1-alpine.

```
student@lfs458-node-1a0a:~$ kubectl edit ds ds-two
....
- image: nginx:1.8.1-alpine
.....
```

16. View the age of the DaemonSets. It should be around ten minutes old, depending on how fast you type.

```
student@lfs458-node-1a0a:~$ kubectl get ds ds-two
NAME      DESIRED    CURRENT    READY    UP-TO-DATE    AVAILABLE    NODE-SELECTOR    AGE
ds-two    2          2          2        2             2           <none>           10m
```

17. Now view the age of the Pods. Two should be much younger than the DaemonSet. They are also a few seconds apart due to the nature of the rolling update where one then the other pod was terminated and recreated.

```
student@lfs458-node-1a0a:~$ kubectl get po
NAME                READY    STATUS    RESTARTS   AGE
ds-one-qc72k        1/1      Running   0           36m
ds-one-z31r4        1/1      Running   0           1h
ds-two-2p8vz        1/1      Running   0           34s
ds-two-8lx7k        1/1      Running   0           32s
```

18. Verify the Pods are using the new version of the software.

```
student@lfs458-node-1a0a:~$ kubectl describe po ds-two-8lx7k |grep Image:
Image:                nginx:1.8.1-alpine
```

19. View the rollout status and the history of the DaemonSets.

```
student@lfs458-node-1a0a:~$ kubectl rollout status ds ds-two
daemon set "ds-two" successfully rolled out
```

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds ds-two
daemonsets "ds-two"
REVISION    CHANGE-CAUSE
1           <none>
2           <none>
```

20. View the changes in the update they should look the same as the previous history, but did not require the Pods to be deleted for the update to take place.

```
student@lfs458-node-1a0a:~$ kubectl rollout history ds ds-two --revision=2
...
Image:                nginx:1.8.1-alpine
```

21. Clean up the system by removing one of the DaemonSets. We will leave the other running.

```
student@lfs458-node-1a0a:~$ kubectl delete ds ds-two
daemonset.extensions "ds-two" deleted
```


Chapter 8

Services



8.1 Labs

Exercise 8.1: Deploy A New Service

Overview

Services (also called **microservices**) are objects which declare a policy to access a logical set of Pods. They are typically assigned with `labels` to allow persistent access to a resource, when front or back end containers are terminated and replaced.

Native applications can use the `Endpoints` API for access. Non-native applications can use a Virtual IP-based bridge to access back end pods. `ServiceTypes` Type could be:

- **ClusterIP** default - exposes on a cluster-internal IP. Only reachable within cluster
- **NodePort** Exposes node IP at a static port. A `ClusterIP` is also automatically created.
- **LoadBalancer** Exposes service externally using cloud providers load balancer. `NodePort` and `ClusterIP` automatically created.
- **ExternalName** Maps service to contents of `externalName` using a `CNAME` record.

We use services as part of decoupling such that any agent or object can be replaced without interruption to access from client to back end application.

Deploy A New Service

1. Deploy two **nginx** servers using **kubectl** and a new `.yaml` file. We will use the `v1beta` version of the API. The kind should be `Deployment` and label it with `nginx`. Create two replicas and expose port 8080. What follows is a well documented file. There is no need to include the comments when you create the file. This file can also be found among the other examples in the tarball.

```
student@lfs458-node-1a0a:~$ vim nginx-one.yaml
```

```

apiVersion: extensions/v1beta1
# Determines YAML versioned schema.
kind: Deployment
# Describes the resource defined in this file.
metadata:
  name: nginx-one
  labels:
    system: secondary
# Required string which defines object within namespace.
  namespace: accounting
# Existing namespace resource will be deployed into.
spec:
  replicas: 2
# How many Pods of following containers to deploy
  template:
    metadata:
      labels:
        app: nginx
# Some string meaningful to users, not cluster. Keys
# must be unique for each object. Allows for mapping
# to customer needs.
    spec:
      containers:
# Array of objects describing containerized application with a Pod.
# Referenced with shorthand spec.template.spec.containers
        - image: nginx:1.7.9
# The Docker image to deploy
          imagePullPolicy: Always
          name: nginx
# Unique name for each container, use local or Docker repo image
          ports:
            - containerPort: 8080
              protocol: TCP
# Optional resources this container may need to function.
          nodeSelector:
            system: secondOne
# One method of node affinity.

```

2. View the existing labels on the nodes in the cluster.

```

student@lfs458-node-1a0a:~$ kubectl get nodes --show-labels
<output_omitted>

```

3. Run the following command and look for the errors. Assuming there is no typo, you should have gotten an error about the accounting namespace.

```

student@lfs458-node-1a0a:~$ kubectl create -f nginx-one.yaml
Error from server (NotFound): error when creating
"nginx-one.yaml": namespaces "accounting" not found

```

4. Create the namespace and try to create the deployment again. There should be no errors this time.

```

student@lfs458-node-1a0a:~$ kubectl create ns accounting
namespace/accounting" created

student@lfs458-node-1a0a:~$ kubectl create -f nginx-one.yaml
deployment.extensions/nginx-one created

```

5. View the status of the new nodes. Note they do not show a Running status.

```

student@lfs458-node-1a0a:~$ kubectl -n accounting get pods
NAME                                READY    STATUS    RESTARTS   AGE
nginx-one-74dd9d578d-fcpmv          0/1     Pending   0           4m
nginx-one-74dd9d578d-r2d67          0/1     Pending   0           4m

```

6. View the node each has been assigned to (or not) and the reason, which shows under events at the end of the output.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting describe pod \
    nginx-one-74dd9d578d-fcpmv

Name:          nginx-one-74dd9d578d-fcpmv
Namespace:     accounting
Node:          <none>

<output_omitted>

Events:
  Type            Reason              Age             From            ....
  ----            -
Warning          FailedScheduling    37s (x25 over 2m29s)  default-scheduler
0/2 nodes are available: 2 node(s) didn't match node selector.
```

7. Label the secondary node. Verify the labels.

```
student@lfs458-node-1a0a:~$ kubectl label node lfs458-worker \
    system=secondOne
node/lfs458-worker labeled

student@lfs458-node-1a0a:~$ kubectl get nodes --show-labels
NAME                STATUS    ROLES    AGE      VERSION    LABELS
lfs458-node-1a0a    Ready    master   1d1h     v1.12.1    \
    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/
hostname=lfs458-node-1a0a,node-role.kubernetes.io/master=
lfs458-worker       Ready    <none>   1d1h     v1.12.1    \
    beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/
hostname=lfs458-worker,system=secondOne
```

8. View the pods in the accounting namespace. They may still show as Pending. Depending on how long it has been since you attempted deployment the system may not have checked for the label. If the Pods show Pending after a minute delete one of the pods. They should both show as Running after a deletion. A change in state will cause the Deployment controller to check the status of both Pods.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting get pods
NAME                READY    STATUS    RESTARTS   AGE
nginx-one-74dd9d578d-fcpmv  1/1      Running    0           10m
nginx-one-74dd9d578d-sts5l  1/1      Running    0           3s
```

9. View Pods by the label we set in the YAML file. If you look back the Pods were given a label of app=nginx.

```
student@lfs458-node-1a0a:~$ kubectl get pods -l app=nginx --all-namespaces
NAMESPACE   NAME                READY    STATUS    RESTARTS   AGE
accounting   nginx-one-74dd9d578d-fcpmv  1/1      Running    0           20m
accounting   nginx-one-74dd9d578d-sts5l  1/1      Running    0           9m
```

10. Recall that we exposed port 8080 in the YAML file. Expose the new deployment.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting expose deployment nginx-one
service/nginx-one exposed
```

11. View the newly exposed endpoints. Note that port 8080 has been exposed on each Pod.

```
student@lfs458-node-9q6r:~$ kubectl -n accounting get ep nginx-one
NAME      ENDPOINTS                                AGE
nginx-one  192.168.1.72:8080,192.168.1.73:8080    47s
```

12. Attempt to access the Pod on port 8080, then on port 80. Even though we exposed port 8080 of the container the application within has not been configured to listen on this port. The **nginx** server will listen on port 80 by default. A `curl` command to that port should return the typical welcome page.

```
student@lfs458-node-1a0a:~$ curl 192.168.1.72:8080
curl: (7) Failed to connect to 192.168.1.72 port 8080: Connection refused

student@lfs458-node-1a0a:~$ curl 192.168.1.72:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

13. Delete the deployment. Edit the YAML file to expose port 80 and create the deployment again.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting delete deploy nginx-one
deployment.extensions "nginx-one" deleted

student@lfs458-node-1a0a:~$ vim nginx-one.yaml

student@lfs458-node-1a0a:~$ kubectl create -f nginx-one.yaml
deployment.extensions/nginx-one created
```

Exercise 8.2: Configure a NodePort

In a previous exercise we deployed a `LoadBalancer` which deployed a `ClusterIP` and `NodePort` automatically. In this exercise we will deploy a `NodePort`. While you can access a container from within the cluster, one can use a `NodePort` to NAT traffic from outside the cluster. One reason to deploy a `NodePort` instead, is that a `LoadBalancer` is also a load balancer resource from cloud providers like GKE and AWS.

1. In a previous step we were able to view the **nginx** page using the internal Pod IP address. Now expose the deployment using the `--type=NodePort`. We will also give it an easy to remember name and place it in the `accounting` namespace. We could pass the port as well, which could help with opening ports in the firewall.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting expose deployment \
    nginx-one --type=NodePort --name=service-lab
service/service-lab exposed
```

2. View the details of the services in the `accounting` namespace. We are looking for the autogenerated port.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting describe services
....
NodePort:                <unset>    32103/TCP
....
```

3. Locate the exterior facing IP address of the cluster. As we are using GCP nodes, which we access via a `FloatingIP`, we will first check the internal only public IP address. Look for the Kubernetes master URL.

```
student@lfs458-node-1a0a:~$ kubectl cluster-info
Kubernetes master is running at https://10.128.0.3:6443
KubeDNS is running at https://10.128.0.3:6443/api/v1/namespaces/
kube-system/services/kube-dns/proxy
To further debug and diagnose cluster problems, use
'kubectl cluster-info dump'.
```

4. Test access to the **nginx** web server using the combination of master URL and `NodePort`.

```
student@lfs458-node-1a0a:~$ curl http://10.128.0.3:32103
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
```

- Using the browser on your local system, use the public IP address you use to SSH into your node and the port. You should still see the **nginx** default page.

Exercise 8.3: Use Labels to Manage Resources

- Try to delete all Pods with the `app=nginx` label, in all namespaces. You should receive an error as this function must be narrowed to a particular namespace. Then delete using the appropriate namespace.

```
student@lfs458-node-1a0a:~$ kubectl delete pods -l app=nginx \
--all-namespaces
Error: unknown flag: --all-namespaces
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl -n accounting delete pods -l app=nginx
pod "nginx-one-74dd9d578d-fcpmv" deleted
pod "nginx-one-74dd9d578d-sts5l" deleted
```

- View the Pods again. New versions of the Pods should be running as the controller responsible for them continues.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-one-74dd9d578d-ddt5r          1/1     Running   0           1m
nginx-one-74dd9d578d-hfzml          1/1     Running   0           1m
```

- We also gave a label to the deployment. View the deployment in the accounting namespace.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting get deploy --show-labels
NAME          DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE  LABELS
nginx-one 2        2        2            2          27m  system=secondary
```

- Delete the deployment using its label.

```
student@lfs458-node-1a0a:~$ kubectl -n accounting delete deploy \
-l system=secondary
deployment.extensions/nginx-one deleted
```

- Remove the label from the secondary node. Note that the syntax is a minus sign directly after the key you want to remove, or `system` in this case.

```
student@lfs458-node-1a0a:~$ kubectl label node lfs458-worker system-
node/lfs458-worker labeled
```


Chapter 9

Volumes and Data



9.1 Labs

Exercise 9.1: Create a ConfigMap

Overview

Container files are ephemeral, which can be problematic for some applications. Should a container be restarted the files will be lost. In addition, we need a method to share files between containers inside a Pod.

A `Volume` is a directory accessible to containers in a Pod. Cloud providers offer volumes which persist further than the life of the Pod, such that AWS or GCE volumes could be pre-populated and offered to Pods, or transferred from one Pod to another. **Ceph** is also another popular solution for dynamic, persistent volumes.

Unlike current **Docker** volumes a Kubernetes volume has the lifetime of the Pod, not the containers within. You can also use different types of volumes in the same Pod simultaneously, but Volumes cannot mount in a nested fashion. Each must have their own mount point. Volumes are declared with `spec.volumes` and mount points with `spec.containers.volumeMounts` parameters. Each particular volume type, 24 currently, may have other restrictions. <https://kubernetes.io/docs/concepts/storage/volumes/#types-of-volumes>

We will also work with a `ConfigMap`, which is basically a set of key-value pairs. This data can be made available so that a Pod can read the data as environment variables or configuration data. A `ConfigMap` is similar to a `Secret`, except they are not base64 byte encoded arrays. They are stored as strings and can be read in serialized form.

Create a ConfigMap

There are three different ways a `ConfigMap` can ingest data, from a literal value, from a file or from a directory of files.

1. We will create a `ConfigMap` containing primary colors. We will create a series of files to ingest into the `ConfigMap`. First, we create a directory `primary` and populate it with four files. Then we create a file in our home directory with our favorite color.

```

student@lfs458-node-1a0a:~$ mkdir primary

student@lfs458-node-1a0a:~$ echo c > primary/cyan

student@lfs458-node-1a0a:~$ echo m > primary/magenta

student@lfs458-node-1a0a:~$ echo y > primary/yellow

student@lfs458-node-1a0a:~$ echo k > primary/black

student@lfs458-node-1a0a:~$ echo "known as key" >> primary/black

student@lfs458-node-1a0a:~$ echo blue > favorite

```

2. Now we will create the ConfigMap and populate it with the files we created as well as a literal value from the command line.

```

student@lfs458-node-1a0a:~$ kubectl create configmap colors \
    --from-literal=text=black \
    --from-file=./favorite \
    --from-file=./primary/
configmap/colors created

```

3. View how the data is organized inside the cluster.

```

student@lfs458-node-1a0a:~$ kubectl get configmap colors
NAME      DATA      AGE
colors    6          30s

student@lfs458-node-1a0a:~$ kubectl get configmap colors -o yaml
apiVersion: v1
data:
  black: |
    k
    known as key
  cyan: |
    c
  favorite: |
    blue
  magenta: |
    m
  text: black
  yellow: |
    y
kind: ConfigMap
<output_omitted>

```

4. Now we can create a Pod to use the ConfigMap. In this case a particular parameter is being defined as an environment variable.

```

student@lfs458-node-1a0a:~$ vim simpleshell.yaml

apiVersion: v1
kind: Pod
metadata:
  name: shell-demo
spec:
  containers:
  - name: nginx
    image: nginx
    env:

```

```
- name: ilike
  valueFrom:
    configMapKeyRef:
      name: colors
      key: favorite
```

5. Create the Pod and view the environmental variable. After you view the parameter, exit out and delete the pod.

```
student@lfs458-node-1a0a:~$ kubectl create -f simpleshell.yaml
pod/shell-demo created

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo \
    -- /bin/bash -c 'echo $ilike'
blue

student@lfs458-node-1a0a:~$ kubectl delete pod shell-demo
pod "shell-demo" deleted
```

6. All variables from a file can be included as environment variables as well. Comment out the previous `env:` stanza and add a slightly different `envFrom` to the file. Having new and old code at the same time can be helpful to see and understand the differences. Recreate the Pod, check all variables and delete the pod again. They can be found spread throughout the environment variable output.

```
student@lfs458-node-1a0a:~$ vim simpleshell.yaml
<output_omitted>
  image: nginx
#   env:
#   - name: ilike
#     valueFrom:
#       configMapKeyRef:
#         name: colors
#         key: favorite
  envFrom:
    - configMapRef:
      name: colors

student@lfs458-node-1a0a:~$ kubectl create -f simpleshell.yaml
pod/shell-demo created

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo \
    -- /bin/bash -c 'env'
HOSTNAME=shell-demo
NJS_VERSION=1.13.6.0.1.14-1~stretch
NGINX_VERSION=1.13.6-1~stretch
black=k
know as key

favorite=blue
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl delete pod shell-demo
pod "shell-demo" deleted
```

7. A ConfigMap can also be created from a YAML file. Create one with a few parameters to describe a car.

```
student@lfs458-node-1a0a:~$ vim car-map.yaml

apiVersion: v1
kind: ConfigMap
metadata:
  name: fast-car
```

```

    namespace: default
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby

```

8. Create the ConfigMap and verify the settings.

```

student@lfs458-node-1a0a:~$ kubectl create -f car-map.yaml
configmap/fast-car created

student@lfs458-node-1a0a:~$ kubectl get configmap fast-car -o yaml
apiVersion: v1
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby
kind: ConfigMap
<output_omitted>

```

9. We will now make the ConfigMap available to a Pod as a mounted volume. You can again comment out the previous environmental settings and add the following new stanza. The containers: and volumes: entries are indented the same number of spaces.

```

student@lfs458-node-1a0a:~$ vim simpleshell.yaml
<output_omitted>
spec:
  containers:
    - name: nginx
      image: nginx
      volumeMounts:
        - name: car-vol
          mountPath: /etc/cars
  volumes:
    - name: car-vol
      configMap:
        name: fast-car
<comment out rest of file>

```

10. Create the Pod again. Verify the volume exists and the contents of a file within. Due to the lack of a carriage return in the file your next prompt may be on the same line as the output, Shelby.

```

student@lfs458-node-1a0a:~$ kubectl create -f simpleshell.yaml
pod "shell-demo" created

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo -- \
    /bin/bash -c 'df -ha |grep car'
/dev/sda1          20G  4.7G   15G   25% /etc/cars

student@lfs458-node-1a0a:~$ kubectl exec -it shell-demo -- \
    /bin/bash -c 'cat /etc/cars/car.trim'
Shelby

```

11. Delete the Pod and ConfigMaps we were using.

```

student@lfs458-node-1a0a:~$ kubectl delete pods shell-demo
pod "shell-demo" deleted

student@lfs458-node-1a0a:~$ kubectl delete configmap fast-car colors
configmap "fast-car" deleted
configmap "colors" deleted

```

Exercise 9.2: Creating a Persistent NFS Volume (PV)

We will first deploy an NFS server. Once tested we will create a persistent NFS volume for containers to claim.

1. Install the software on your master node.

```
student@lfs458-node-1a0a:~$ sudo apt-get update && sudo \
    apt-get install -y nfs-kernel-server
<output_omitted>
```

2. Make and populate a directory to be shared. Also give it similar permissions to `/tmp/`

```
student@lfs458-node-1a0a:~$ sudo mkdir /opt/sfw

student@lfs458-node-1a0a:~$ sudo chmod 1777 /opt/sfw/

student@lfs458-node-1a0a:~$ sudo bash -c \
    'echo software > /opt/sfw/hello.txt'
```

3. Edit the NFS server file to share out the newly created directory. In this case we will share the directory with all. You can always **snoop** to see the inbound request in a later step and update the file to be more narrow.

```
student@lfs458-node-1a0a:~$ sudo vim /etc/exports
/opt/sfw/ *(rw, sync, no_root_squash, subtree_check)
```

4. Cause `/etc/exports` to be re-read:

```
student@lfs458-node-1a0a:~$ sudo exportfs -ra
```

5. Test by mounting the resource from your **second** node.

```
student@lfs458-worker:~$ sudo apt-get -y install nfs-common
<output_omitted>

student@lfs458-worker:~$ showmount -e lfs458-node-1a0a
Export list for lfs458-node-1a0a:
/opt/sfw *
```

```
student@lfs458-worker:~$ sudo mount 10.128.0.3:/opt/sfw /mnt

student@lfs458-worker:~$ ls -l /mnt
total 4
-rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt
```

6. Return to the master node and create a YAML file for the object with kind, `PersistentVolume`. Use the hostname of the master server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the resource will not start. Note that the `accessModes` do not currently affect actual access and are typically used as labels instead.

```
student@lfs458-node-1a0a:~$ vim PVol.yaml

apiVersion: v1
kind: PersistentVolume
metadata:
  name: pvvol-1
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
```

```
nfs:
  path: /opt/sfw
  server: lfs458-node-1a0a  #<-- Edit to match master node
  readOnly: false
```

7. Create the persistent volume, then verify its creation.

```
student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml
persistentvolume/pvvol-1 created

student@lfs458-node-1a0a:~$ kubectl get pv
NAME          CAPACITY  ACCESSMODES  RECLAIMPOLICY  STATUS
CLAIM         STORAGECLASS  REASON    AGE
pvvol-1       1Gi         RWX         Retain         Available  4s
```

Exercise 9.3: Creating a Persistent Volume Claim (PVC)

Before Pods can take advantage of the new PV we need to create a **Persistent Volume Claim (PVC)**.

1. Begin by determining if any currently exist.

```
student@lfs458-node-1a0a:~$ kubectl get pvc
No resources found.
```

2. Create a YAML file for the new pvc.

```
student@lfs458-node-1a0a:~$ vim pvc.yaml
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-one
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 200Mi
```

3. Create and verify the new pvc is bound. Note that the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used.

```
student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml
persistentvolumeclaim/pvc-one created
```

```
student@lfs458-node-1a0a:~$ kubectl get pvc
NAME      STATUS  VOLUME  CAPACITY  ACCESSMODES  STORAGECLASS  AGE
pvc-one   Bound   pvvol-1  1Gi       RWX           default/pvc-one  4s
```

4. Look at the status of the pv again, to determine if it is in use. It should show a status of Bound.

```
student@lfs458-node-1a0a:~$ kubectl get pv
NAME          CAPACITY  ACCESSMODES  RECLAIMPOLICY  STATUS  CLAIM          STORAGECLASS  REASON  AGE
pvvol-1       1Gi         RWX         Retain         Bound   default/pvc-one  5m
```

5. Create a new deployment to use the pvc. We will copy and edit an existing deployment yaml file. We will change the deployment name then add a volumeMounts section under containers and volumes section to the general spec. The name used must match in both places, whatever name you use. The claimName must match an existing pvc. As shown in the following example.

```
student@lfs458-node-1a0a:~$ cp first.yaml nfs-pod.yaml
```

```
student@lfs458-node-1a0a:~$ vim nfs-pod.yaml
```

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  generation: 1
  labels:
    run: nginx
  name: nginx-nfs
  namespace: default
  resourceVersion: "1411"
spec:
  replicas: 1
  selector:
    matchLabels:
      run: nginx
  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
    type: RollingUpdate
  template:
    metadata:
      creationTimestamp: null
      labels:
        run: nginx
    spec:
      containers:
      - image: nginx
        imagePullPolicy: Always
        name: nginx
        volumeMounts:
        - name: nfs-vol
          mountPath: /opt
        ports:
        - containerPort: 80
          protocol: TCP
        resources: {}
        terminationMessagePath: /dev/termination-log
        terminationMessagePolicy: File
      volumes:
        #<<-- These four lines
        - name: nfs-vol
          persistentVolumeClaim:
            claimName: pvc-one
      dnsPolicy: ClusterFirst
      restartPolicy: Always
      schedulerName: default-scheduler
      securityContext: {}
      terminationGracePeriodSeconds: 30
```

6. Create the pod using the newly edited file.

```
student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml
deployment.apps/nginx-nfs created
```

7. Look at the details of the pod. You may see the daemonset pods running as well.

```
student@lfs458-node-1a0a:~$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
nginx-nfs-1054709768-s8g28         1/1      Running   0           3m

student@lfs458-node-1a0a:~$ kubectl describe pod nginx-nfs-1054709768-s8g28
Name:                                nginx-nfs-1054709768-s8g28
Namespace:                          default
Node:                                lfs458-worker/10.128.0.5

<output_omitted>

Mounts:
  /opt from nfs-vol (rw)

<output_omitted>

Volumes:
  nfs-vol:
    Type:          PersistentVolumeClaim (a reference to a PersistentV...
    ClaimName:      pvc-one
    ReadOnly:       false
<output_omitted>
```

8. View the status of the PVC. It should show as bound.

```
student@lfs458-node-1a0a:~$ kubectl get pvc
NAME      STATUS VOLUME  CAPACITY  ACCESS MODES  STORAGECLASS  AGE
pvc-one   Bound  pvvol-1  1Gi       RWX            pvvol-1       2m
```

Exercise 9.4: Using a ResourceQuota to Limit PVC Count and Usage

The flexibility of cloud-based storage often requires limiting consumption among users. We will use the ResourceQuota object to both limit the total consumption as well as the number of persistent volume claims.

1. Begin by deleting the deployment we had created to use NFS, the pv and the pvc.

```
student@lfs458-node-1a0a:~$ kubectl delete deploy nginx-nfs
deployment.extensions "nginx-nfs" deleted

student@lfs458-node-1a0a:~$ kubectl delete pvc pvc-one
persistentvolumeclaim "pvc-one" deleted

student@lfs458-node-1a0a:~$ kubectl delete pv pvvol-1
persistentvolume "pvvol-1" deleted
```

2. Create a yaml file for the ResourceQuota object. Set the storage limit to ten claims with a total usage of 500Mi.

```
student@lfs458-node-1a0a:~$ vim storage-quota.yaml

apiVersion: v1
kind: ResourceQuota
metadata:
  name: storagequota
spec:
  hard:
    persistentvolumeclaims: "10"
    requests.storage: "500Mi"
```


3. Create a new namespace called `small`. View the namespace information prior to the new quota. Either the long name with double dashes `--namespace` or the nickname `ns` work for the resource.

```
student@lfs458-node-1a0a:~$ kubectl create namespace small
namespace/small created
```

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
Name:          small
Labels:        <none>
Annotations:    <none>
Status:        Active
```

No resource quota.

No resource limits.

4. Create a new pv and pvc in the `small` namespace.

```
student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml -n small
persistentvolume/pvvol-1 created
```

```
student@lfs458-node-1a0a:~$ kubectl create -f pvc.yaml -n small
persistentvolumeclaim/pvc-one created
```

5. Create the new resource quota, placing this object into the `low-usage-limit` namespace.

```
student@lfs458-node-1a0a:~$ kubectl create -f storage-quota.yaml \
-n small
resourcequota/storagequota created
```

6. Verify the `small` namespace has quotas. Compare the output to the same command above.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
Name:          small
Labels:        <none>
Annotations:    <none>
Status:        Active
```

Resource Quotas

Name:	storagequota
Resource	Used Hard
-----	---
persistentvolumeclaims	1 10
requests.storage	200Mi 500Mi

No resource limits.

7. Remove the namespace line from the `nfs-pod.yaml` file. Should be around line 11 or so. This will allow us to pass other namespaces on the command line.

```
student@lfs458-node-1a0a:~$ vim nfs-pod.yaml
```

8. Create the container again.

```
student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml \
-n small
deployment.apps/nginx-nfs created
```

9. Determine if the deployment has a running pod.

```
student@lfs458-node-1a0a:~$ kubectl get deploy --namespace=small
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx-nfs     1         1         1             0           43s

student@lfs458-node-1a0a:~$ kubectl -n small describe deploy \
    nginx-nfs
<output_omitted>
```

10. Look to see if the pods are ready.

```
student@lfs458-node-1a0a:~$ kubectl get po -n small
NAME                                READY   STATUS    RESTARTS   AGE
nginx-nfs-2854978848-g3khf         1/1     Running   0          37s
```

11. Ensure the Pod is running and is using the NFS mounted volume. If you pass the namespace first Tab will auto-complete the pod name.

```
student@lfs458-node-1a0a:~$ kubectl -n small describe po \
    nginx-nfs-2854978848-g3khf
Name:          nginx-nfs-2854978848-g3khf
Namespace:     small
<output_omitted>

Mounts:
  /opt from nfs-vol (rw)
<output_omitted>
```

12. View the quota usage of the namespace

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
```

```
Resource Quotas
Name:          storagequota
Resource      Used      Hard
-----
persistentvolumeclaims 1        10
requests.storage 200Mi    500Mi
```

No resource limits.

13. Create a 300M file inside of the `/opt/sfw` directory on the host and view the quota usage again. Note that with NFS the size of the share is not counted against the deployment.

```
student@lfs458-node-1a0a:~$ sudo dd if=/dev/zero \
    of=/opt/sfw/bigfile bs=1M count=300
300+0 records in
300+0 records out
314572800 bytes (315 MB, 300 MiB) copied, 0.196794 s, 1.6 GB/s
```

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
```

```
Resource Quotas
Name:          storagequota
Resource      Used      Hard
-----
persistentvolumeclaims 1        10
requests.storage 200Mi    500Mi
```

<output_omitted>

```
student@lfs458-node-1a0a:~$ du -h /opt/
301M    /opt/sfw
```

```
41M      /opt/cni/bin
41M      /opt/cni
341M     /opt/
```

14. Now let us illustrate what happens when a deployment requests more than the quota. Begin by shutting down the existing deployment.

```
student@lfs458-node-1a0a:~$ kubectl -n small get deploy
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx-nfs     1         1         1             1           11m

student@lfs458-node-1a0a:~$ kubectl -n small delete deploy nginx-nfs
deployment.extensions "nginx-nfs" deleted
```

15. Once the Pod has shut down view the resource usage of the namespace again. Note the storage did not get cleaned up when the pod was shut down.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
Resource Quotas
Name:
Resource          Used      Hard
-----
persistentvolumeclaims 1         10
requests.storage      200Mi     500Mi
```

16. Remove the pvc then view the pv it was using. Note the RECLAIM POLICY and STATUS.

```
student@lfs458-node-1a0a:~$ kubectl get pvc -n small
NAME      STATUS    VOLUME   CAPACITY   ACCESSMODES   STORAGECLASS   AGE
pvc-one   Bound     pvvol-1  1Gi        RWX            small           19m

student@lfs458-node-1a0a:~$ kubectl -n small delete pvc pvc-one
persistentvolumeclaim "pvc-one" deleted

student@lfs458-node-1a0a:~$ kubectl -n small get pv
NAME      CAPACITY   ACCESSMODES   RECLAIMPOLICY   STATUS   CLAIM
STORAGECLASS  REASON    AGE
pvvol-1  1Gi    RWX    Retain   Released   small/pvc-one  44m
```

17. Dynamically provisioned storage uses the ReclaimPolicy of the StorageClass which could be Delete, Retain, or some types allow Recycle. Manually created persistent volumes default to Retain unless set otherwise at creation. The default storage policy is to retain the storage to allow recovery of any data. To change this begin by viewing the yaml output.

```
student@lfs458-node-1a0a:~$ kubectl get pv/pvvol-1 -o yaml
....
  path: /opt/sfw
  server: lfs458-node-1a0a
  persistentVolumeReclaimPolicy: Retain
status:
  phase: Released
```

18. Currently we will need to delete and re-create the object. Future development on a deleter plugin is planned. We will re-create the volume and allow it to use the Retain policy, then change it once running.

```
student@lfs458-node-1a0a:~$ kubectl delete pv/pvvol-1
persistentvolume "pvvol-1" deleted

student@lfs458-node-1a0a:~$ grep Retain PVol.yaml
  persistentVolumeReclaimPolicy: Retain

student@lfs458-node-1a0a:~$ kubectl create -f PVol.yaml
persistentvolume "pvvol-1" created
```

19. We will use `kubectl patch` to change the retention policy to `Delete`. The yaml output from before can be helpful in getting the correct syntax.

```
student@lfs458-node-1a0a:~$ kubectl patch pv pvvol-1 -p \
'{"spec":{"persistentVolumeReclaimPolicy":"Delete"}}'
```

persistentvolume/pvvol-1 patched

```
student@lfs458-node-1a0a:~$ kubectl get pv/pvvol-1
```

NAME	CAPACITY	ACCESSMODES	RECLAIMPOLICY	STATUS	CLAIM
pvvol-1	1Gi	RWX	Delete	Available	2m

20. View the current quota settings.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
```

```
.
requests.storage      0      500Mi
```

21. Create the pvc again. Even with no pods running, note the resource usage.

```
student@lfs458-node-1a0a:~$ kubectl -n small create -f pvc.yaml
```

persistentvolumeclaim/pvc-one created

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
```

```
.
requests.storage      200Mi      500Mi
```

22. Remove the existing quota from the namespace.

```
student@lfs458-node-1a0a:~$ kubectl -n small get resourcequota
```

NAME	CREATED AT
storagequota	2018-08-01T04:10:02Z

```
student@lfs458-node-1a0a:~$ kubectl -n small delete \
resourcequota storagequota
```

resourcequota "storagequota" deleted

23. Edit the `storagequota.yaml` file and lower the capacity to 100Mi.

```
student@lfs458-node-1a0a:~$ vim storage-quota.yaml
```

```
.
requests.storage: "100Mi"
```

24. Create and verify the new storage quota. Note the hard limit has already been exceeded.

```
student@lfs458-node-1a0a:~$ kubectl create -f storage-quota.yaml -n small
```

resourcequota/storagequota created

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
```

```
.
persistentvolumeclaims      1      10
requests.storage      200Mi      100Mi
```

No resource limits.

25. Create the deployment again. View the deployment. Note there are no errors seen.

```
student@lfs458-node-1a0a:~$ kubectl create -f nfs-pod.yaml \
-n small
deployment.apps/nginx-nfs created

student@lfs458-node-1a0a:~$ kubectl -n small describe deploy/nginx-nfs
Name:                nginx-nfs
Namespace:           small
<output_omitted>
```

26. Examine the pods to see if they are actually running.

```
student@lfs458-node-1a0a:~$ kubectl -n small get po
NAME                                READY   STATUS    RESTARTS   AGE
nginx-nfs-2854978848-vb6bh         1/1     Running   0           58s
```

27. As we were able to deploy more pods even with apparent hard quota set, let us test to see if the reclaim of storage takes place. Remove the deployment and the persistent volume claim.

```
student@lfs458-node-1a0a:~$ kubectl -n small delete deploy nginx-nfs
deployment.extensions "nginx-nfs" deleted

student@lfs458-node-1a0a:~$ kubectl -n small delete pvc/pvc-one
persistentvolumeclaim "pvc-one" deleted
```

28. View if the persistent volume exists. You will see it attempted a removal, but failed. If you look closer you will find the error has to do with the lack of a deleter volume plugin for NFS. Other storage protocols have a plugin.

```
student@lfs458-node-1a0a:~$ kubectl -n small get pv
NAME      CAPACITY   ACCESSMODES   RECLAIMPOLICY   STATUS   CLAIM
STORAGECLASS   REASON   AGE
pvvol-1    1Gi       RWX          Delete         Failed   small/pvc-one  20m
```

29. Ensure the deployment, pvc and pv are all removed.

```
student@lfs458-node-1a0a:~$ kubectl delete pv/pvvol-1
persistentvolume "pvvol-1" deleted
```

30. Edit the persistent volume YAML file and change the persistentVolumeReclaimPolicy: to Recycle.

```
student@lfs458-node-1a0a:~$ vim PVol.yaml
....
  persistentVolumeReclaimPolicy: Recycle
....
```

31. Add a LimitRange to the namespace and attempt to create the persistent volume and persistent volume claim again. We can use the LimitRange we used earlier.

```
student@lfs458-node-1a0a:~$ kubectl -n small create -f \
  low-resource-range.yaml
limitrange/low-resource-range created
```

32. View the settings for the namespace. Both quotas and resource limits should be seen.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
Resource Limits
-----
Type      Resource  Min  Max  Default Request  Default Limit  ...
-----
Container cpu      -    -    500m              1              -
Container memory   -    -    100Mi             500Mi          -
```

33. Create the persistent volume again. View the resource. Note the Reclaim Policy is Recycle.

```
student@lfs458-node-1a0a:~$ kubectl -n small create -f PVol.yaml
persistentvolume/pvvol-1 created

student@lfs458-node-1a0a:~$ kubectl get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  ...
pvvol-1       1Gi       RWX           Recycle         Available  ...
```

34. Attempt to create the persistent volume claim again. The quota only takes effect if there is also a resource limit in effect.

```
student@lfs458-node-1a0a:~$ kubectl -n small create -f pvc.yaml
Error from server (Forbidden): error when creating "pvc.yaml":
  persistentvolumeclaims "pvc-one" is forbidden: exceeded quota:
storagequota, requested: requests.storage=200Mi, used:
requests.storage=0, limited: requests.storage=100Mi
```

35. Edit the resourcequota to increase the requests.storage to 500mi.

```
student@lfs458-node-1a0a:~$ kubectl -n small edit resourcequota
....
spec:
  hard:
    persistentvolumeclaims: "10"
    requests.storage: 500Mi
status:
  hard:
    persistentvolumeclaims: "10"
....
```

36. Create the pvc again. It should work this time. Then create the deployment again.

```
student@lfs458-node-1a0a:~$ kubectl -n small create -f pvc.yaml
persistentvolumeclaim/pvc-one created

student@lfs458-node-1a0a:~$ kubectl -n small create -f nfs-pod.yaml
deployment.apps/nginx-nfs created
```

37. View the namespace settings.

```
student@lfs458-node-1a0a:~$ kubectl describe ns small
<output_omitted>
```

38. Delete the deployment. View the status of the pv and pvc.

```
student@lfs458-node-1a0a:~$ kubectl -n small delete deploy nginx-nfs
deployment.extensions "nginx-nfs" deleted

student@lfs458-node-1a0a:~$ kubectl get pvc -n small
NAME      STATUS  VOLUME  CAPACITY  ACCESS MODES  STORAGECLASS  AGE
pvc-one   Bound   pvvol-1  1Gi       RWX           Recycle       7m

student@lfs458-node-1a0a:~$ kubectl -n small get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM  STORA...
pvvol-1       1Gi       RWX           Recycle         Bound   small/pvc-one  ...
```

39. Delete the pvc and check the status of the pv. It should show as Available.

```
student@lfs458-node-1a0a:~$ kubectl -n small delete pvc pvc-one
persistentvolumeclaim "pvc-one" deleted

student@lfs458-node-1a0a:~$ kubectl -n small get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM  STORA...
pvvol-1       1Gi       RWX           Recycle         Available  ...
```

40. Remove the pv and any other resources created during this lab.

```
student@lfs458-node-1a0a:~$ kubectl delete pv pvvol-1  
persistentvolume "pvvol-1" deleted
```


Chapter 10

Ingress



10.1 Labs

Exercise 10.1: Advanced Service Exposure

Configure an Ingress Controller

With such a fast changing project, it is important to keep track of updates. The main place to find documentation of the current version is <https://kubernetes.io/>.

1. If you have a large number of services to expose outside of the cluster, or to expose a low-number port on the host node you can deploy an ingress controller or a service mesh. While **nginx** and **GCE** have controllers officially supported by Kubernetes.io, the **Traefik** ingress controller is easier to install. At the moment.

```
student@lfs458-node-1a0a:~$ kubectl create deployment secondapp \
  --image=nginx
```

2. Find the labels currently in use by the deployment. We will use them to tie traffic from the ingress controller to the proper service.

```
student@lfs458-node-1a0a:~$ kubectl get deployments secondapp -o yaml |grep label -A2
labels:
  app: secondapp
  name: secondapp
--
      labels:
        app: secondapp
      spec:
```

3. Expose the new server as a NodePort.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment secondapp \
  --type=NodePort --port=80
```

4. As we have RBAC configured we need to make sure the controller will run and be able to work with all necessary ports, endpoints and resources. Create a YAML file to declare a clusterrole and a clusterrolebinding.

```
student@lfs458-node-1a0a:~$ vim ingress.rbac.yaml
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: traefik-ingress-controller
rules:
  - apiGroups:
    - ""
    resources:
    - services
    - endpoints
    - secrets
    verbs:
    - get
    - list
    - watch
  - apiGroups:
    - extensions
    resources:
    - ingresses
    verbs:
    - get
    - list
    - watch
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: traefik-ingress-controller
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: traefik-ingress-controller
subjects:
- kind: ServiceAccount
  name: traefik-ingress-controller
  namespace: kube-system
```

5. Create the new role and binding.

```
student@lfs458-node-1a0a:~$ kubectl create -f ingress.rbac.yaml
clusterrole.rbac.authorization.k8s.io "traefik-ingress-controller" created
clusterrolebinding.rbac.authorization.k8s.io "traefik-ingress-controller" created
```

6. Create the Traefik controller. We will use a script directly from their website. This URL has a shorter version below:
<https://raw.githubusercontent.com/containous/traefik/master/examples/k8s/traefik-ds.yaml>

```
student@lfs458-node-1a0a:~$ wget https://tinyurl.com/yawpexdt -O traefik-ds.yaml
```

7. We need to take out some security context settings, such that the diff output between the new and old would be true. Add the hostNetwork line and remove the securityContext lines. The indentation for hostNetwork should line up with the containers: line.

```
student@lfs458-node-1a0a:~$ vim traefik-ds.yaml
diff traefik-ds.yaml.1 ds/traefik-ds.yaml
23a24      ## Add this line
>      hostNetwork: true
34,39d34      ## Remove these lines
<      securityContext:
<      capabilities:
```

```
<      drop:
<      - ALL
<      add:
<      - NET_BIND_SERVICE
```

8. Then create the ingress controller using **kubectl create**

```
student@lfs458-node-1a0a:~$ kubectl create -f traefik-ds.yaml
serviceaccount "traefik-ingress-controller" created
daemonset.extensions "traefik-ingress-controller" created
service "traefik-ingress-service" created
```

9. Now that there is a new controller we need to pass some rules, so it knows how to handle requests. Note that the host mentioned is `www.example.com`, which is probably not your node name. We will pass a false header when testing. Also the service name needs to match the `secondapp` label we found in an earlier step.

```
student@lfs458-node-1a0a:~$ vim ingress.rule.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-test
  annotations:
    kubernetes.io/ingress.class: traefik
spec:
  rules:
  - host: www.example.com
    http:
      paths:
      - backend:
          serviceName: secondapp
          servicePort: 80
        path: /
```

10. Now ingest the rule into the cluster.

```
student@lfs458-node-1a0a:~$ kubectl create -f ingress.rule.yaml
ingress.extensions "ingress-test" created
```

11. We should be able to test the internal and external IP addresses, and see the nginx welcome page. The loadbalancer would present the traffic, a **curl** request in this case, to the externally facing interface. Use **ip a** to find the IP address of the interface which would face the loadbalancer. In this example the interface would be `ens4`, and the IP would be `10.128.0.7`.

```
student@lfs458-node-1a0a:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens4: <BROADCAST,UP,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group default qlen 1000
    link/ether 42:01:0a:80:00:03 brd ff:ff:ff:ff:ff:ff
    inet 10.128.0.7/32 brd 10.128.0.3 scope global ens4
        valid_lft forever preferred_lft forever
<output_omitted>

student@lfs458-node-1a0a:~$ curl -H "Host: www.example.com" http://10.128.0.7/
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

```
student@lfs458-node-1a0a:~$ curl -H "Host: www.example.com" http://35.193.3.179
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>
```

12. At this point we would keep adding more and more web servers. We'll configure one more, which would then be a process continued as many times as desired.

Begin by deploying another nginx server. Give it a label and expose port 80.

```
student@lfs458-node-1a0a:~$ kubectl create deployment thirdpage --image=nginx
deployment.apps/thirdpage created
```

13. Find the label for the new deployment. Look for the `name:`, which would be `thirdpage` in this example.

```
student@lfs458-node-1a0a:~$ kubectl get deployment thirdpage -o yaml |grep -A2 Label
labels:
  app: thirdpage
  name: thirdpage
--
      labels:
        app: thirdpage
      spec:
```

14. Expose the new server as a NodePort.

```
student@lfs458-node-1a0a:~$ kubectl expose deployment \
  thirdpage --type=NodePort --port=80
service/thirdpage exposed
```

15. Now we will customize the installation. Run a bash shell inside the new pod. Your pod name will end differently. Install **vim** inside the container then edit the `index.html` file of nginx so that the title of the web page will be Third Page.

```
student@lfs458-node-1a0a:~$ kubectl exec -it thirdpage-5cf8d67664-zcmfh -- /bin/bash

root@thirdpage-5cf8d67664-zcmfh:/# apt-get update
<output_omitted>

root@thirdpage-5cf8d67664-zcmfh:/# apt-get install vim -y
<output_omitted>

root@thirdpage-5cf8d67664-zcmfh:/# vim /usr/share/nginx/html/index.html
<!DOCTYPE html>
<html>
<head>
<title>Third Page</title>
<style>
<output_omitted>
```

16. Edit the ingress rules to point the `thirdpage` service. Use the `serviceName` we found in an earlier step of `thirdpage`.

```
student@lfs458-node-1a0a:~$ kubectl edit ingress ingress-test
<output_omitted>
- host: www.example.com
  http:
    paths:
    - backend:
        serviceName: secondapp
        servicePort: 80
      path: /
```

```
- host: thirdpage.org
  http:
    paths:
      - backend:
          serviceName: thirdpage
          servicePort: 80
        path: /
    status:
<output_omitted>
```

17. Test the second hostname using **curl** locally as well as from a remote system.

```
student@lfs458-node-1a0a:~$ curl -H "Host: thirdpage.org" http://10.128.0.7/
<!DOCTYPE html>
<html>
<head>
<title>Third Page</title>
<style>
<output_omitted>
```


Chapter 11

Scheduling



11.1 Labs

Exercise 11.1: Assign Pods Using Labels

Overview

While allowing the system to distribute Pods on your behalf is typically the best route, you may want to determine which nodes a Pod will use. For example you may have particular hardware requirements to meet for the workload. You may want to assign VIP Pods to new, faster hardware and everyone else to older hardware.

In this exercise we will use `labels` to schedule Pods to a particular node. Then we will explore `taints` to have more flexible deployment in a large environment.

Assign Pods Using Labels

1. Begin by getting a list of the nodes. They should be in the ready state and without added labels or taints.

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE    VERSION
lfs458-node-1a0a    Ready    master   2d     v1.12.1
lfs458-worker       Ready    <none>   2d     v1.12.1
```

2. View the current labels and taints for the nodes.

```
student@lfs458-node-1a0a:~$ kubectl describe nodes |grep -i label
Labels:                beta.kubernetes.io/arch=amd64
Labels:                beta.kubernetes.io/arch=amd64

student@lfs458-node-1a0a:~$ kubectl describe nodes |grep -i taint
Taints:                <none>
Taints:                <none>
```

3. Verify there are no deployments running, outside of the `kube-system` namespace. If there are, delete them. Then get a count of how many containers are running on both the master and secondary nodes. There are about 24 containers running on the master in the following example, and eight running on the worker. There are status lines which increase the **wc** count. You may have more or less, depending on previous labs and cleaning up of resources.

```
student@lfs458-node-1a0a:~$ kubectl get deployments --all-namespaces
NAMESPACE      NAME           DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
default        secondapp      1         1         1            1           37m
default        thirdpage      1         1         1            1           14m
kube-system     calico-typha   0         0         0            0           2d15h
kube-system     coredns        2         2         2            2           2d15h
low-usage-limit limited-hog     1         1         1            1           1d29m
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
24
```

```
student@lfs458-worker:~$ sudo docker ps |wc -l
14
```

4. For the purpose of the exercise we will assign the master node to be VIP hardware and the secondary node to be for others.

```
student@lfs458-node-1a0a:~$ kubectl label nodes lfs458-node-1a0a status=vip
node/lfs458-node-1a0a labeled
```

```
student@lfs458-node-1a0a:~$ kubectl label nodes lfs458-worker status=other
node/lfs458-worker labeled
```

5. Verify your settings. You will also find there are some built in labels such as `hostname`, `os` and `architecture` type. The output below appears on multiple lines for readability.

```
student@lfs458-node-1a0a:~$ kubectl get nodes --show-labels
NAME              STATUS    ROLES    AGE   VERSION   LABELS
lfs458-node-1a0a  Ready    master   2d    v1.12.1   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=lfs458-node-1a0a,node-role.kubernetes.io/master=,status=vip
lfs458-worker     Ready    <none>   2d    v1.12.1   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=lfs458-worker,status=other
```

6. Create `vip.yaml` to spawn four busybox containers which sleep the whole time. Include the `nodeSelector` entry.

```
student@lfs458-node-1a0a:~$ vim vip.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: vip
spec:
  containers:
  - name: vip1
    image: busybox
    args:
    - sleep
    - "1000000"
  - name: vip2
    image: busybox
    args:
    - sleep
    - "1000000"
  - name: vip3
```



```

    image: busybox
    args:
    - sleep
    - "1000000"
- name: vip4
  image: busybox
  args:
  - sleep
  - "1000000"
nodeSelector:
  status: vip

```

7. Deploy the new pod. Verify the containers have been created on the master node. It may take a few seconds for all the containers to spawn. Check both the master and the secondary nodes.

```

student@lfs458-node-1a0a:~$ kubectl create -f vip.yaml
pod/vip created

```

```

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
29

```

```

student@lfs458-worker:~$ sudo docker ps |wc -l
8

```

8. Delete the pod then edit the file, commenting out the nodeSelector lines. It may take a while for the containers to fully terminate.

```

student@lfs458-node-1a0a:~$ kubectl delete pod vip
pod "vip" deleted

```

```

student@lfs458-node-1a0a:~$ vim vip.yaml
....
# nodeSelector:
#   status: vip

```

9. Create the pod again. Containers should now be spawning on both nodes. You may see pods for the daemonsets as well.

```

student@lfs458-node-1a0a:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
ds-one-bdqst                        1/1     Running   0           145m
ds-one-t2t7z                        1/1     Running   0           158m
secondapp-85765cd95c-2q9sx          1/1     Running   0           43m
thirdpage-7c9b56bfd-2q5pr           1/1     Running   0           20m

student@lfs458-node-1a0a:~$ kubectl create -f vip.yaml
pod/vip created

```

10. Determine where the new containers have been deployed. They should be more evenly spread this time.

```

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
24

```

```

student@lfs458-worker:~$ sudo docker ps |wc -l
19

```

11. Create another file for other users. Change the names from vip to others, and uncomment the nodeSelector lines.

```

student@lfs458-node-1a0a:~$ cp vip.yaml other.yaml

student@lfs458-node-1a0a:~$ sed -i s/vip/other/g other.yaml

```

```
student@lfs458-node-1a0a:~$ vim other.yaml
.
nodeSelector:
  status: other
```

12. Create the other containers. Determine where they deploy.

```
student@lfs458-node-1a0a:~$ kubectl create -f other.yaml
pod/other created

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
24

student@lfs458-worker:~$ sudo docker ps |wc -l
24
```

13. Shut down both pods and verify they terminated. Only our previous pods should be found.

```
student@lfs458-node-1a0a:~$ kubectl delete pods vip other
pod "vip" deleted
pod "other" deleted

student@lfs458-node-1a0a:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
ds-one-bdqst                        1/1     Running   0           153m
ds-one-t2t7z                        1/1     Running   0           166m
secondapp-85765cd95c-2q9sx          1/1     Running   0           51m
thirdpage-7c9b56bfdd-2q5pr          1/1     Running   0           28m
```

Exercise 11.2: Using Taints to Control Pod Deployment

Use taints to manage where Pods are deployed or allowed to run. In addition to assigning a Pod to a group of nodes, you may also want to limit usage on a node or fully evacuate Pods. Using taints is one way to achieve this. You may remember that the master node begins with a `NoSchedule` taint. We will work with three taints to limit or remove running pods.

1. Verify that the master and secondary node have the minimal number of containers running.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment secondapp \
thirdpage
deployment.extensions "secondapp" deleted
deployment.extensions "thirdpage" deleted
```

2. Create a deployment which will deploy eight **nginx** containers. Begin by creating a YAML file.

```
student@lfs458-node-1a0a:~$ vim taint.yaml

apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: taint-deployment
spec:
  replicas: 8
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
```

```
image: nginx:1.7.9
ports:
- containerPort: 80
```

3. Apply the file to create the deployment.

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment.apps/taint-deployment created
```

4. Determine where the containers are running. In the following example three have been deployed on the master node and five on the secondary node. Remember there will be other housekeeping containers created as well. Your numbers may be slightly different.

```
student@lfs458-node-1a0a:~$ sudo docker ps |grep nginx
00c1be5df1e7      nginx@sha256:e3456c851a152494c3e..
<output_omitted>
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
28
student@lfs458-worker:~$ sudo docker ps |wc -l
26
```

5. Delete the deployment. Verify the containers are gone.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment.extensions "taint-deployment" deleted
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
24
```

6. Now we will use a taint to affect the deployment of new containers. There are three taints, NoSchedule, PreferNoSchedule and NoExecute. The taints having to do with schedules will be used to determine newly deployed containers, but will not affect running containers. The use of NoExecute will cause running containers to move.

Taint the secondary node, verify it has the taint then create the deployment again. We will use the key of bubba to illustrate the key name is just some string an admin can use to track Pods.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-worker \
      bubba=value:PreferNoSchedule
node/lfs458-worker tainted
```

```
student@lfs458-node-1a0a:~$ kubectl describe node |grep Taint
Taints:      bubba=value:PreferNoSchedule
Taints:      <none>
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment.apps/taint-deployment created
```

7. Locate where the containers are running. We can see that more containers are on the master, but there still were some created on the secondary. Delete the deployment when you have gathered the numbers.

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
32
```

```
student@lfs458-worker:~$ sudo docker ps |wc -l
22
```

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment.extensions "taint-deployment" deleted
```

8. Remove the taint, verify it has been removed. Note that the key is used with a minus sign appended to the end.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-worker bubba-
node/lfs458-worker untainted
```

```
student@lfs458-node-1a0a:~$ kubectl describe node |grep Taint
Taints:                <none>
Taints:                <none>
```

9. This time use the NoSchedule taint, then create the deployment again. The secondary node should not have any new containers, with only daemonsets and other essential pods running.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-worker \
    bubba=value:NoSchedule
node/lfs458-worker tainted
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment.apps/taint-deployment created
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
24
```

```
student@lfs458-worker:~$ sudo docker ps |wc -l
14
```

10. Remove the taint and delete the deployment. When you have determined that all the containers are terminated create the deployment again. Without any taint the containers should be spread across both nodes.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment.extensions "taint-deployment" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-worker bubba-
node/lfs458-worker untainted
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment.apps/taint-deployment created
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
32
```

```
student@lfs458-worker:~$ sudo docker ps |wc -l
22
```

11. Now use the NoExecute to taint the secondary node. Wait a minute then determine if the containers have moved. The DNS containers can take a while to shutdown. A few containers will remain on the worker node to continue communication from the cluster.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-worker \
    bubba=value:NoExecute
node "lfs458-worker" tainted
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
32
```

```
student@lfs458-worker:~$ sudo docker ps |wc -l
6
```

12. Remove the taint. Wait a minute. Note that all of the containers did not return to their previous placement.

```
student@lfs458-node-1a0a:~$ kubectl taint nodes lfs458-worker bubba-
node/lfs458-worker untainted
```

```
student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
32
```

```
student@lfs458-worker:~$ sudo docker ps |wc -l
6
```

13. In addition to the ability to taint a node you can also set the status to drain. First view the status, then destroy the existing deployment. Note that the status reports Ready, even though it will not allow containers to be executed. Also note that the output mentioned that DaemonSet-managed pods are not affected by default, as we saw in an earlier lab. This time lets take a closer look at what happens to existing pods and nodes.

Existing containers are not moved, but no new containers are created. You may receive an error
error: unable to drain node "<your node>", aborting command...

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE    VERSION
lfs458-node-1a0a    Ready    master   2d     v1.12.1
lfs458-worker       Ready    <none>   2d     v1.12.1

student@lfs458-node-1a0a:~$ kubectl drain lfs458-worker
node/lfs458-worker cordoned
error: DaemonSet-managed pods (use --ignore-daemonsets to ignore): kube-flannel-ds-fx3tx, kube-proxy-q2q4k
```

14. Verify the state change of the node. It should indicate no new Pods will be scheduled.

```
student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS              ROLES    AGE    VERSION
lfs458-node-1a0a    Ready              master   2d     v1.12.1
lfs458-worker       Ready,SchedulingDisabled <none>   2d     v1.12.1
```

15. Delete the deployment to destroy the current Pods.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment.extensions "taint-deployment" deleted
```

16. Create the deployment again and determine where the containers have been deployed.

```
student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment.apps/taint-deployment created

student@lfs458-node-1a0a:~$ sudo docker ps |wc -l
44
```

17. Return the status to Ready, then destroy and create the deployment again. The containers should be spread across the nodes. Begin by removing the cordon on the node.

```
student@lfs458-node-1a0a:~$ kubectl uncordon lfs458-worker
node/lfs458-worker uncordoned

student@lfs458-node-1a0a:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE    VERSION
lfs458-node-1a0a    Ready    master   2d     v1.12.1
lfs458-worker       Ready    <none>   2d     v1.12.1
```

18. Delete and re-create the deployment.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment
deployment.extensions "taint-deployment" deleted

student@lfs458-node-1a0a:~$ kubectl apply -f taint.yaml
deployment.apps/taint-deployment created
```

19. View the **docker ps** output again. Both nodes should have almost the same number of containers deployed. The master will have a few more, due to its role.
20. Remove the deployment a final time to free up resources.

```
student@lfs458-node-1a0a:~$ kubectl delete deployment taint-deployment  
deployment.extensions "taint-deployment" deleted
```

Chapter 12

Logging and Troubleshooting



12.1 Labs

Exercise 12.1: Review Log File Locations

Overview

In addition to various logs files and command output, you can use **journalctl** to view logs from the node perspective. We will view common locations of log files, then a command to view container logs. There are other logging options, such as the use of a **sidecar** container dedicated to loading the logs of another container in a pod.

Whole cluster logging is not yet available with Kubernetes. Outside software is typically used, such as **Fluentd**, part of <https://fluentd.org/>, which is another member project of **CNCF**, like Kubernetes.

Review Log File Locations

Take a quick look at the following log files and web sites. As server processes move from node level to running in containers the logging also moves.

1. If using a **systemd** based Kubernetes cluster view the node level logs for **kubelet**, the local Kubernetes agent. Each node will have different contents as this is node specific.

```
student@lfs458-node-1a0a:~$ journalctl -u kubelet |less
<output_omitted>
```

2. Major Kubernetes processes now run in containers. You can view them from the container or the pod perspective. Use the **find** command to locate the **kube-apiserver** log. Your output will be different, but will be very long. Once you locate the files use the **diff** utility to compare them. There should be no difference, as they are symbolic links to [/var/log/pods/](#). If you follow the links the log files are unique.

```
student@lfs458-node-1a0a:~$ sudo find / -name "*apiserver*log"
/var/log/containers/kube-apiserver-u16-12-1-dcb8_kube-system_kube-apiserver-
eddae7079382cd382cd55f8f46b192565dd16b6858206039d49b1ad4693c2a10.log
/var/log/containers/kube-apiserver-u16-12-1-dcb8_kube-system_kube-apiserver-
d00a48877af4ed4c7f8eedf2c7805c77cfabb31fcb453f7d89ffa52fc6ea5f36.log

student@lfs458-node-1a0a:~$ sudo diff /var/log/containers/kube-apiserver-u16-
12-1-dcb8_kube-system_kube-apiserver-eddae7079382cd382cd55f8f46b192565dd16b68
58206039d49b1ad4693c2a10.log /var/log/containers/kube-apiserver-u16-12-1-
dcb8_kube-system_kube-apiserver-d00a48877af4ed4c7f8eedf2c7805c77cfabb31fcb453
f7d89ffa52fc6ea5f36.log
<output_omitted>
```

3. Take a look at the log file.

```
student@lfs458-node-1a0a:~$ sudo less /var/log/containers/kube-apiserver-u16-
12-1-dcb8_kube-system_kube-apiserver-d00a48877af4ed4c7f8eedf2c7805c77cfabb31f
cb453f7d89ffa52fc6ea5f36.log
```

4. Search for and review other log files for kube-dns, kube-flannel, and kube-proxy.
5. If not on a Kubernetes cluster using **systemd** you can view the text files on the master node.

- (a) `/var/log/kube-apiserver.log`
Responsible for serving the API
- (b) `/var/log/kube-scheduler.log`
Responsible for making scheduling decisions
- (c) `/var/log/kube-controller-manager.log`
Controller that manages replication controllers

6. `/var/log/containers`
Various container logs

7. `/var/log/pods/`
More log files for current Pods.

8. Worker Nodes Files (on non-**systemd** systems)

- (a) `/var/log/kubelet.log`
Responsible for running containers on the node
- (b) `/var/log/kube-proxy.log`
Responsible for service load balancing

9. More reading: <https://kubernetes.io/docs/tasks/debug-application-cluster/\debug-service/> and <https://kubernetes.io/docs/tasks/debug-application-cluster/\determine-reason-pod-failure/>

Exercise 12.2: Viewing Logs Output

Container standard out can be seen via the **kubectl logs** command. If there is no standard out, you would not see any output. In addition, the logs would be destroyed if the container is destroyed.

1. View the current Pods in the cluster. Be sure to view Pods in all namespaces.

```
student@lfs458-node-1a0a:~$ kubectl get po --all-namespaces
NAMESPACE      NAME                                READY   STATUS    RESTARTS   AGE
default        ds-one-qc72k                       1/1     Running   0           3h
default        ds-one-z31r4                       1/1     Running   0           3h
....
kube-system    etcd-lfs458-node-1a0a              1/1     Running   2           9h
```



```

kube-system    kube-apiserver-lfs458-node-1a0a    1/1    Running    2    9h
kube-system    kube-controller-manager-lfs458-node-1a0a    1/1    Running    2    9h
kube-system    kube-dns-2425271678-w80vx    3/3    Running    6    9h
kube-system    kube-scheduler-lfs458-node-1a0a    1/1    Running    2    9h
...

```

2. View the logs associated with various infrastructure pods. Using the **Tab** key you can get a list and choose a container. Then you can start typing the name of a pod and use **Tab** to complete the name.

```

student@lfs458-node-1a0a:~$ kubectl -n kube-system logs <Tab><Tab>
calico-etcd-n6h2q
etcd-lfs458-1-11-1update-cm35
calico-kube-controllers-74b888b647-9ds42
kube-apiserver-lfs458-1-11-1update-cm35
calico-node-6j8hc
kube-controller-manager-lfs458-1-11-1update-cm35
calico-node-dq6kf
kube-proxy-8sn6f
coredns-78fcd6894-7fpfp
kube-proxy-wf5dr
coredns-78fcd6894-g6k99
kube-scheduler-lfs458-1-11-1update-cm35

```

```

student@lfs458-node-1a0a:~$ kubectl -n kube-system logs \
    kube-apiserver-lfs458-1-11-1update-cm35
Flag --insecure-port has been deprecated, This flag will be
removed in a future version.
I0729 21:29:23.026394      1 server.go:703] external host
was not specified, using 10.128.0.2
I0729 21:29:23.026667      1 server.go:145] Version: v1.11.1
I0729 21:29:23.784000      1 plugins.go:158] Loaded 8 mutating
admission controller(s) successfully in the following order:
NamespaceLifecycle,LimitRanger,ServiceAccount,NodeRestriction,
Priority,DefaultTolerationSeconds,DefaultStorageClass,
MutatingAdmissionWebhook.
I0729 21:29:23.784025      1 plugins.go:161] Loaded 6 validating
admission controller(s) successfully in the following order:
LimitRanger,ServiceAccount,Priority,PersistentVolumeClaimResize,
ValidatingAdmissionWebhook,ResourceQuota.
<output_omitted>

```

3. View the logs of other Pods in your cluster.

Chapter 13

Custom Resource Definition



13.1 Labs

Exercise 13.1: Create a Custom Resource Definition

Overview

ThirdPartyResource is no longer included with the API in v1.8 and its use will return a validation error. If you have upgraded from a version prior to Kubernetes v1.7, you will need to convert them to CustomResourceDefinitions (CRD). A new resource often requires a controller to manage the resource. Creation of the controller is beyond the scope of this course, basically it is a watch-loop comparing a spec file to the current state and making changes until the states match. A good discussion of creating a controller can be found here: <https://coreos.com/blog/introducing-operators.html>.

Create a Custom Resource Definition

We will make a simple CRD, but without any particular action. It will be enough to find the object ingested into the API and responding to commands.

1. We will create a new YAML file.

```
student@lfs458-node-1a0a:~$ vim crd.yaml
```

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: crontabs.training.lfs458.com
    # This name must match names below.
    # <plural>.<group> syntax
spec:
  scope: Cluster      #Could also be Namespaced
  group: training.lfs458.com
```

```

version: v1
names:
  kind: CronTab      #Typically CamelCased for resource manifest
  plural: crontabs   #Shown in URL
  singular: crontab  #Short name for CLI alias
  shortNames:
  - ct               #CLI short name

```

2. Add the new resource to the cluster.

```

student@lfs458-node-1a0a:~$ kubectl create -f crd.yaml
customresourcedefinition.apiextensions.k8s.io/crontabs.training.lfs458.com
created

```

3. View and describe the resource. You'll note the **describe** output is unlike other objects we have seen so far.

```

student@lfs458-node-1a0a:~$ kubectl get crd
NAME                                CREATED AT
crontabs.training.lfs458.com        2018-08-03T05:25:20Z
<output_omitted>

student@lfs458-node-1a0a:~$ kubectl describe crd crontab<Tab>
Name:          crontabs.training.lfs458.com
Namespace:
Labels:         <none>
Annotations:    <none>
API Version:    apiextensions.k8s.io/v1beta1
Kind:           CustomResourceDefinition
<output_omitted>

```

4. Now that we have a new API resource we can create a new object of that type. In this case it will be a crontab-like image, which does not actually exist, but is being used for demonstration.

```

student@lfs458-node-1a0a:~$ vim new-crontab.yaml

apiVersion: "training.lfs458.com/v1"
  # This is from the group and version of new CRD
kind: CronTab
  # The kind from the new CRD
metadata:
  name: new-cron-object
spec:
  cronSpec: "*/5 * * * *"
  image: some-cron-image
  #Does not exist

```

5. Create the new object and view the resource using short and long name.

```

student@lfs458-node-1a0a:~$ kubectl create -f new-crontab.yaml
crontab.training.lfs458.com/new-cron-object created

student@lfs458-node-1a0a:~$ kubectl get CronTab
NAME          AGE
new-cron-object 22s

student@lfs458-node-1a0a:~$ kubectl get ct
NAME          AGE
new-cron-object 29s

student@lfs458-node-1a0a:~$ kubectl describe ct

```

```
Name:          new-cron-object
Namespace:
Labels:        <none>
```

```
<output_omitted>
```

```
Spec:
  Cron Spec:  */5 * * * *
  Image:      some-cron-image
  Events:     <none>
```

6. To clean up the resources we will delete the CRD. This should delete all of the endpoints and objects using it as well.

```
student@lfs458-node-1a0a:~$ kubectl delete -f crd.yaml
customresourcedefinition.apiextensions.k8s.io
"crontabs.training.lfs458.com" deleted
```

```
student@lfs458-node-1a0a:~$ kubectl get ct
Error from server (NotFound): Unable to list "crontabs": the server
could not find the requested resource
(get crontabs.training.lfs458.com)
```


Chapter 14

Kubernetes Federation



14.1 Labs

There is no lab to complete for this chapter.

Chapter 15

Helm



15.1 Labs

Exercise 15.1: Working with Helm and Charts

Overview

helm allows for easy deployment of complex configurations. This could be handy for a vendor to deploy a multi-part application in a single step. Through the use of a Chart, or template file, the required components and their relationships are declared. Local agents like **Tiller** use the API to create objects on your behalf. Effectively its orchestration for orchestration.

There are a few ways to install **Helm**. The newest version may require building from source code. We will download a recent, stable version. Once installed we will deploy a Chart, which will configure **Hadoop** on our cluster.

Install Helm

1. On the master node use **wget** to download the compressed tar file. The short URL below is for: <https://storage.googleapis.com/kubernetes-helm/helm-v2.7.0-linux-amd64.tar.gz>

```
student@lfs458-node-1a0a:~$ wget goo.gl/nbEcHn
<output_omitted>
nbEcHn      100%[=====] 11.61M  --.-KB/s   in 0.1s

2018-08-03 05:34:56 (91.7 MB/s) - nbEcHn saved [12169373/12169373]
```

2. Uncompress and expand the file.

```
student@lfs458-node-1a0a:~$ tar -xvf nbEcHn
linux-amd64/
linux-amd64/README.md
linux-amd64/helm
linux-amd64/LICENSE
```

- Copy the **helm** binary to the `/usr/local/bin/` directory, so it is usable via the shell search path.

```
student@lfs458-node-1a0a:~$ sudo cp linux-amd64/helm /usr/local/bin/
```

- Due to new **RBAC** configuration **helm** is unable to run in the default namespace, in this version of Kubernetes. During initialization you could choose to create and declare a new namespace. Other RBAC issues may be encountered even then. In this lab we will create a service account for **tiller**, and give it admin abilities on the cluster. More on **RBAC** in another chapter.

Begin by creating the serviceaccount object.

```
student@lfs458-node-1a0a:~$ kubectl create serviceaccount \
    --namespace kube-system tiller
serviceaccount "tiller" created
```

- Bind the serviceaccount to the admin role called `cluster-admin` inside the `kube-system` namespace.

```
student@lfs458-node-1a0a:~$ kubectl create clusterrolebinding \
    tiller-cluster-rule \
    --clusterrole=cluster-admin \
    --serviceaccount=kube-system:tiller
clusterrolebinding.rbac.authorization.k8s.io/tiller-cluster-rule created
```

- We can now initialize **helm**. This process will also configure **tiller** the client process. There are several possible options to pass such as `nodeAffinity`, a particular version of software, alternate storage backend, and even a dry-run option to generate JSON or YAML output. The output could be edited and ingested into **kubectl**. We will use default values in this case.

```
student@lfs458-node-1a0a:~$ helm init
```

<output_omitted>

- Update the `tiller-deploy` deployment to have the service account.

```
student@lfs458-node-1a0a:~$ kubectl -n kube-system patch deployment \
    tiller-deploy -p \
    '{"spec":{"template":{"spec":{"serviceAccount":"tiller"}}}}'
deployment.extensions/tiller-deploy patched
```

- Verify the **tiller** pod is running. Examine the logs of the pod. Note that each line of log begins with an tag of the component generating the messages, such as `[main]`, `[storage]`, and `[storage]`.

```
student@lfs458-node-1a0a:~$ kubectl get pods --all-namespaces
kube-system    tiller-deploy-84b97f465c-76lvs    1/1    Running    0 30m
```

```
student@lfs458-node-1a0a:~$ kubectl -n kube-system logs \
    tiller-deploy-84b97f465c-76lvs
<output_omitted>
```

- View the available sub-commands for **helm**. As with other Kubernetes tools, expect ongoing change.

```
student@lfs458-node-1a0a:~$ helm help
<output_omitted>
```

- View the current configuration files, archives and plugins for **helm**. Return to this directory after you have worked with a Chart later in the lab.

```
student@lfs458-node-1a0a:~$ helm home
/home/student/.helm

student@lfs458-node-1a0a:~$ ls -R /home/student/.helm/
/home/student/.helm/:
cache  plugins  repository  starters

/home/student/.helm/cache:
archive
<output_omitted>
```

11. Verify **helm** and **tiller** are responding, also check the current version installed.

```
student@lfs458-node-1a0a:~$ helm version
Client: &version.Version{SemVer:"v2.7.0", GitCommit:"08c1144f5...
Server: &version.Version{SemVer:"v2.7.0", GitCommit:"08c1144f5...
```

12. Ensure both are upgraded to the most recent stable version.

```
student@lfs458-node-1a0a:~$ helm init --upgrade
$HELM_HOME has been configured at /home/student/.helm.

Tiller (the Helm server-side component) has been upgraded
to the current version.
Happy Helming!
```

13. A Chart is a collection of containers to deploy an application. There is a collection available on <https://github.com/kubernetes/charts/tree/master/stable>, provided by vendors, or you can make your own. Take a moment and view the current stable Charts. Then search for available stable databases.

```
student@lfs458-node-1a0a:~$ helm search database
NAME                VERSION      DESCRIPTION
stable/cockroachdb  2.0.3       CockroachDB is a scalable, survivable,...
stable/dokuwiki     3.3.0       DokuWiki is a standards-compliant, ...
stable/janusgraph   0.2.0       Open source, scalable graph database.
stable/kubedb       0.1.3       DEPRECATED KubeDB by AppCode - Making...
stable/mariadb      5.2.3       Fast, reliable, scalable, and easy to use...
<output_omitted>
```

14. We will install the **mariadb**. Take a look at install details <https://github.com/kubernetes/charts/tree/master/stable/mariadb#custom-mycnf-configuration> The **--debug** option will create a lot of output. Note the interesting name for the deployment, like **illmannered-salamander**. The output will typically suggest ways to access the software. As well we will indicate that we do not want persistent storage, which would require use to create an available PV.

```
student@lfs458-node-1a0a:~$ helm --debug install stable/mariadb \
--set master.persistence.enabled=false \
--set slave.persistence.enabled=false

[debug] Created tunnel using local port: '38396'

[debug] SERVER: "localhost:38396"

[debug] Original chart version: ""
[debug] Fetched stable/mariadb to /home/student/.helm/cache/archive/mar...

[debug] CHART PATH: /home/student/.helm/cache/archive/mariadb-2.0.1.tgz

NAME:    illmannered-salamander
<output_omitted>
```

15. Using some of the information at the end of the previous command output we will deploy another container and access the database. We begin by getting the root password for `illmannered-salamander-mariadb`. Be aware the output lacks a carriage return, so the next prompt will appear on the same line. We will need the password to access the running MariaDB database.

```
student@lfs458-node-1a0a:~$ kubectl get secret -n default \
    illmannered-salamander-mariadb \
    -o jsonpath="{.data.mariadb-root-password}" \
    | base64 --decode
IFBldzAQfx
```

16. Now we will install another container to act as a client for the database. We will use **apt-get** to install client software.

```
student@lfs458-node-1a0a:~$ kubectl run -i --tty ubuntu \
    --image=ubuntu:16.04 --restart=Never -- bash -il
If you don't see a command prompt, try pressing enter.
root@ubuntu:/#

root@ubuntu:/# apt-get update ; apt-get install -y mariadb-client
<output_omitted>
```

17. Use the client software to access the database. The following command uses the server name and the root password we found in a previous step. Both of yours will be different.

```
root@ubuntu:/# mysql -h illmannered-salamander-mariadb -p
Enter password: IFBldzAQfx
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MariaDB connection id is 153
Server version: 10.1.28-MariaDB Source distribution

Copyright (c) 2000, 2017, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [(none)]> SHOW DATABASES;
+-----+
| Database                |
+-----+
| information_schema      |
| my_database             |
| mysql                  |
| performance_schema     |
| test                    |
+-----+
5 rows in set (0.00 sec)

MariaDB [(none)]>
MariaDB [(none)]> quit
root@ubuntu:/# exit
```

18. View the Chart history on the system. The use of the **-a** option will show all Charts including deleted and failed attempts. The output below shows the current running Chart as well as a previously deleted **hadoop** Chart.

```
student@lfs458-node-1a0a:~$ helm list -a
NAME          REVISION    UPDATED      STATUS  CHART          NAMESPACE
goodly-beetle 1    Wed Nov  8 23:01:24 2017 DELETED  hadoop-1.0.1  default
illmannered-salamander 1    Thu Nov  9 05:00:12 2017 DEPLOYED mariadb-...
```

19. Delete the **mariadb** Chart. No output should happen from the list.

```
student@lfs458-node-1a0a:~$ helm delete illmannered-salamander  
release "illmannered-salamander" deleted
```

```
student@lfs458-node-1a0a:~$ helm list
```

20. Add another repository and view the Charts available.

```
student@lfs458-node-1a0a:~$ helm repo add common \  
    http://storage.googleapis.com/kubernetes-charts  
"common" has been added to your repositories
```

```
student@lfs458-node-1a0a:~$ helm repo list  
NAME          URL  
stable        https://kubernetes-charts.storage.googleapis.com  
local         http://127.0.0.1:8879/charts  
common        http://storage.googleapis.com/kubernetes-charts
```

```
student@lfs458-node-1a0a:~$ helm search  
NAME                                VERSION    DESCRIPTION  
stable/acs-engine-autoscaler        2.1.0      Scales worker nodes within...  
stable/artifactory                   6.2.0      Universal Repository Manag...  
<output_omitted>
```


Chapter 16

Security



16.1 Labs

Exercise 16.1: Working with TLS

Overview

We have learned that the flow of access to a cluster begins with TLS connectivity, then authentication followed by authorization, finally an admission control plug-in allows advanced features prior to the request being fulfilled. The use of `Initializers` allows the flexibility of a shell-script to dynamically modify the request. As security is an important, ongoing concern, there may be multiple configurations used depending on the needs of the cluster.

Every process making API requests to the cluster must authenticate or be treated as an anonymous user.

Working with TLS

While one can have multiple cluster root Certificate Authorities (CA) by default each cluster uses their own, intended for intra-cluster communication. The CA certificate bundle is distributed to each node and as a secret to default service accounts. The **kubelet** is a local agent which ensures local containers are running and healthy.

1. View the **kubelet** on both the master and secondary nodes. The **kube-apiserver** also shows security information such as certificates and authorization mode. As **kubelet** is a **systemd** service we will start looking at that output.

```
student@lfs458-node-1a0a:~$ systemctl status kubelet.service
kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: en
   Drop-In: /etc/systemd/system/kubelet.service.d
           10-kubeadm.conf
<output_omitted>
```

2. If we look at the status output, and follow the cgroup information, which is a long line we where configuration settings are drawn from, we see where the configuration file can be found.

```
CGroup: /system.slice/kubelet.service
19523 /usr/bin/kubelet .... --config=/var/lib/kubelet/config.yaml ..
```

3. Take a look at the settings in the `/var/lib/kubelet/config.yaml` file. Among other information we can see the `/etc/kubernetes/pki/` directory is used for accessing the **kube-apiserver**. Near the end of the output it also sets the directory to find other pod spec files.

```
student@lfs458-node-1a0a:~$ sudo less /var/lib/kubelet/config.yaml
address: 0.0.0.0
apiVersion: kubelet.config.k8s.io/v1beta1
authentication:
  anonymous:
    enabled: false
  webhook:
    cacheTTL: 2m0s
    enabled: true
x509:
  clientCAFile: /etc/kubernetes/pki/ca.crt
```

4. Other agents on the master node interact with the **kube-apiserver**. View the configuration files where these settings are made. This was set in the previous YAML file. Look at one of the files for cert information.

```
student@lfs458-node-1a0a:~$ sudo ls /etc/kubernetes/manifests/
etcd.yaml          kube-controller-manager.yaml
kube-apiserver.yaml kube-scheduler.yaml

student@lfs458-node-1a0a:~$ sudo less \
    /etc/kubernetes/manifests/kube-controller-manager.yaml
<output_omitted>
```

5. The use of tokens has become central to authorizing component communication. The tokens are kept as **secrets**. Take a look at the current secrets in the `kube-system` namespace.

```
student@lfs458-node-1a0a:~$ kubectl -n kube-system get secrets
NAME                                     TYPE
DATA      AGE
attachdetach-controller-token-xqr8n    kubernet.es.io/service-account-token
3         5d
bootstrap-signer-token-xbp6s           kubernet.es.io/service-account-token
3         5d
bootstrap-token-i3r13t                 bootstrap.kubernet.es.io/token
7         5d
<output_omitted>
```

6. Take a closer look at one of the secrets and the token within. The `certificate-controller-token` could be one to look at. The use of the Tab key can help with long names. Long lines have been truncated in the output below.

```
student@lfs458-node-1a0a:~$ kubectl -n kube-system get secrets \
    certificate<Tab> -o yaml
apiVersion: v1
data:
  ca.crt: LS0tLS1CRUdJTi....
  namespace: a3ViZS1zeXNOZW0=
  token: ZX1KaGJHY2lPaUpTVXpJM....
kind: Secret
metadata:
  annotations:
    kubernet.es.io/service-account.name: certificate-controller
    kubernet.es.io/service-account.uid: 7dfa2aa0-9376-11e8-8cfb
-42010a800002
  creationTimestamp: 2018-07-29T21:29:36Z
  name: certificate-controller-token-wnrwh
```



```

namespace: kube-system
resourceVersion: "196"
selfLink: /api/v1/namespaces/kube-system/secrets/certificate-
controller-token-wnrwh
uid: 7dfbb237-9376-11e8-8cfb-42010a800002
type: kubernetes.io/service-account-token

```

7. The **kubectl config** command can also be used to view and update parameters. When making updates this could avoid a typo removing access to the cluster. View the current configuration settings. The keys and certs are redacted from the output automatically.

```

student@lfs458-node-1a0a:~$ kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
<output_omitted>

```

8. View the options, such as setting a password for the admin instead of a key. Read through the examples and options.

```

student@lfs458-node-1a0a:~$ kubectl config set-credentials -h
Sets a user entry in kubeconfig
<output_omitted>

```

9. Make a copy of your access configuration file. Later steps will update this file and we can view the differences.

```

student@lfs458-node-1a0a:~$ cp ~/.kube/config ~/cluster-api-config

```

10. Explore working with cluster and security configurations both using **kubectl** and **kubeadm**. Among other values, find the name of your cluster. You will need to become root to work with **kubeadm**.

```

student@lfs458-node-1a0a:~$ kubectl config <Tab><Tab>
current-context  get-contexts      set-context      view
delete-cluster  rename-context    set-credentials
delete-context   set               unset
get-clusters     set-cluster       use-context

```

```

student@lfs458-node-1a0a:~$ sudo -i

```

```

root@lfs458-node-1a0a:~# kubeadm token -h
<output_omitted>

```

```

root@lfs458-node-1a0a:~# kubeadm config -h
<output_omitted>

```

11. Review the cluster default configuration settings. At over 150 lines there may be some interesting tidbits to the security and infrastructure of the cluster.

```

student@lfs458-node-1a0a:~$ kubeadm config print-default
api:
  advertiseAddress: 10.128.0.2
  bindPort: 6443
  controlPlaneEndpoint: ""
apiVersion: kubeadm.k8s.io/v1alpha2
auditPolicy:
  logDir: /var/log/kubernetes/audit
  logMaxAge: 2
  path: ""
bootstrapTokens:
- groups:
  - system:bootstrappers:kubeadm:default-node-token
  token: abcdef.0123456789abcdef
<output_omitted>

```

Exercise 16.2: Authentication and Authorization

Kubernetes clusters have two types of users: service accounts and normal users, but normal users are assumed to be managed by an outside service. There are no objects to represent them and they cannot be added via an API call, but service accounts can be added.

We will use **RBAC** to configure access to actions within a namespace for a new contractor, Developer Dan who will be working on a new project.

1. Create two namespaces, one for production and the other for development.

```
student@lfs458-node-1a0a:~$ kubectl create ns development
namespace "development" created
```

```
student@lfs458-node-1a0a:~$ kubectl create ns production
namespace "production" created
```

2. View the current clusters and context available. The context allows you to configure the cluster to use, namespace and user for **kubectl** commands in an easy and consistent manner.

```
student@lfs458-node-1a0a:~$ kubectl config get-contexts
CURRENT  NAME                  CLUSTER          AUTHINFO          NAMESPACE
*         kubernetes-admin@kubernetes  kubernetes       kubernetes-admin
```

3. Create a new user DevDan and assign a password of lfs458.

```
student@lfs458-node-1a0a:~$ sudo useradd -s /bin/bash DevDan
student@lfs458-node-1a0a:~$ sudo passwd DevDan
Enter new UNIX password: lfs458
Retype new UNIX password: lfs458
passwd: password updated successfully
```

4. Generate a private key then Certificate Signing Request (CSR) for DevDan.

```
student@lfs458-node-1a0a:~$ openssl genrsa -out DevDan.key 2048
Generating RSA private key, 2048 bit long modulus
.....+++
.....+++
e is 65537 (0x10001)
```

```
student@lfs458-node-1a0a:~$ openssl req -new -key DevDan.key \
-out DevDan.csr -subj "/CN=DevDan/O=development"
```

5. Using the newly created request generate a self-signed certificate using the x509 protocol. Use the CA keys for the Kubernetes cluster and set a 45 day expiration. You'll need to use **sudo** to access to the inbound files.

```
student@lfs458-node-1a0a:~$ sudo openssl x509 -req -in DevDan.csr \
-CA /etc/kubernetes/pki/ca.crt \
-CAkey /etc/kubernetes/pki/ca.key \
-CAcreateserial \
-out DevDan.crt -days 45
Signature ok
subject=/CN=DevDan/O=development
Getting CA Private Key
```

6. Update the access config file to reference the new key and certificate. Normally we would move them to a safe directory instead of a non-root user's home.

```
student@lfs458-node-1a0a:~$ kubectl config set-credentials DevDan \
--client-certificate=/home/student/DevDan.crt \
--client-key=/home/student/DevDan.key
User "DevDan" set.
```

7. View the update to your credentials file. Use **diff** to compare against the copy we made earlier.

```
student@lfs458-node-1a0a:~$ diff cluster-api-config .kube/config
9a10,14
>     namespace: development
>     user: DevDan
>     name: DevDan-context
> - context:
>     cluster: kubernetes
15a21,25
> - name: DevDan
>     user:
>     as-user-extra: {}
>     client-certificate: /home/student/DevDan.crt
>     client-key: /home/student/DevDan.key
```

8. We will now create a context. For this we will need the name of the cluster, namespace and CN of the user we set or saw in previous steps.

```
student@lfs458-node-1a0a:~$ kubectl config set-context DevDan-context \
    --cluster=kubernetes \
    --namespace=development \
    --user=DevDan
Context "DevDan-context" created.
```

9. Attempt to view the Pods inside the DevDan-context. Be aware you will get an error.

```
student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context get pods
Error from server (Forbidden): pods is forbidden: User "DevDan"
cannot list pods in the namespace "development"
```

10. Verify the context has been properly set.

```
student@lfs458-node-1a0a:~$ kubectl config get-contexts
CURRENT  NAME                CLUSTER      AUTHINFO      NAMESPACE
*        DevDan-context      kubernetes   DevDan        development
*        kubernetes-admin@kubernetes kubernetes   kubernetes-admin
```

11. Again check the recent changes to the cluster access config file.

```
student@lfs458-node-1a0a:~$ diff cluster-api-config .kube/config
9a10,14
>     namespace: development
>     user: DevDan
>     name: DevDan-context
> - context:
>     cluster: kubernetes
15a21,25
> - name: DevDan
>     user:
<output_omitted>
```

12. We will now create a YAML file to associate RBAC rights to a particular namespace and Role.

```
student@lfs458-node-1a0a:~$ vim role-dev.yaml
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: development
  name: developer
rules:
- apiGroups: ["", "extensions", "apps"]
  resources: ["deployments", "replicasets", "pods"]
  verbs: ["list", "get", "watch", "create", "update", "patch", "delete"]
# You can use ["*"] for all verbs
```

13. Create the object. Check white space and for typos if you encounter errors.

```
student@lfs458-node-1a0a:~$ kubectl create -f role-dev.yaml
role.rbac.authorization.k8s.io/developer created
```

14. Now we create a RoleBinding to associate the Role we just created with a user. Create the object when the file has been created.

```
student@lfs458-node-1a0a:~$ vim rolebind.yaml
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: developer-role-binding
  namespace: development
subjects:
- kind: User
  name: DevDan
  apiGroup: ""
roleRef:
  kind: Role
  name: developer
  apiGroup: ""
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f rolebind.yaml
rolebinding.rbac.authorization.k8s.io/developer-role-binding created
```

15. Test the context again. This time it should work. There are no Pods running so you should get a response of No resources found.

```
student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context get pods
No resources found.
```

16. Create a new pod, verify it exists, then delete it.

```
student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context \
  create deployment nginx --image=nginx
deployment.apps/nginx created
```

```
student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-7c87f569d-7gb9k              1/1     Running   0           5s
```

```
student@lfs458-node-1a0a:~$ kubectl --context=DevDan-context delete deploy nginx
deployment.extensions "nginx" deleted
```

17. We will now create a different context for production systems. The Role will only have the ability to view, but not create or delete resources. Begin by copying and editing the Role and RoleBindings YAML files.

```
student@lfs458-node-1a0a:~$ cp role-dev.yaml role-prod.yaml
```

```
student@lfs458-node-1a0a:~$ vim role-prod.yaml
```

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: production      #<<- This line
  name: dev-prod             #<<- and this line
rules:
- apiGroups: ["", "extensions", "apps"]
  resources: ["deployments", "replicasets", "pods"]
  verbs: ["get", "list", "watch"] #<<- and this one
```

```
student@lfs458-node-1a0a:~$ cp rolebind.yaml rolebindprod.yaml
```

```
student@lfs458-node-1a0a:~$ vim rolebindprod.yaml
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: production-role-binding
  namespace: production
subjects:
- kind: User
  name: DevDan
  apiGroup: ""
roleRef:
  kind: Role
  name: dev-prod
  apiGroup: ""
```

18. Create both new objects.

```
student@lfs458-node-1a0a:~$ kubectl apply -f role-prod.yaml
role.rbac.authorization.k8s.io/dev-prod created
```

```
student@lfs458-node-1a0a:~$ kubectl apply -f rolebindprod.yaml
rolebinding.rbac.authorization.k8s.io/production-role-binding created
```

19. Create the new context for production use.

```
student@lfs458-node-1a0a:~$ kubectl config set-context ProdDan-context \
  --cluster=kubernetes \
  --namespace=production \
  --user=DevDan
Context "ProdDan-context" created.
```

20. Verify that user DevDan can view pods using the new context.

```
student@lfs458-node-1a0a:~$ kubectl --context=ProdDan-context get pods
No resources found.
```

21. Try to create a Pod in production. The developer should be Forbidden.

```
student@lfs458-node-1a0a:~$ kubectl --context=ProdDan-context run \
  nginx --image=nginx
Error from server (Forbidden): deployments.extensions is forbidden: User "DevDan" cannot \
  create deployments.extensions in the namespace "production"
```

22. View the details of a role.

```
student@lfs458-node-1a0a:~$ kubectl describe role dev-prod -n production
Name:          dev-prod
Labels:        <none>
Annotations:   kubectl.kubernetes.io/last-applied-configuration=
{"apiVersion":"rbac.authorization.k8s.io/v1beta1","kind":"Role",
,"metadata":{"annotations":{},"name":"dev-prod","namespace":
"production"},"rules":[{"api...
PolicyRule:
  Resources          Non-Resource URLs  Resource Names      Verbs
  -----
  deployments        []                 []                  [get list watch]
  deployments.apps   []                 []                  [get list watch]
<output_omitted>
```

23. Experiment with other subcommands in both contexts. They should match those listed in the respective roles.

Exercise 16.3: Admission Controllers

The last stop before a request is sent to the API server is an admission control plug-in. They interact with features such as setting parameters like a default storage class, checking resource quotas, or security settings. A newer feature (v1.7.x) is dynamic controllers which allow new controllers to be ingested or configured at runtime.

1. View the current admission controller settings. Unlike earlier versions of Kubernetes the controllers are now compiled into the server, instead of being passed at run-time. Instead of a list of which controllers to use we can enable and disable specific plugins.

```
student@lfs458-node-1a0a:~$ sudo grep admission \  
    /etc/kubernetes/manifests/kube-apiserver.yaml  
- --disable-admission-plugins=PersistentVolumeLabel  
- --enable-admission-plugins=NodeRestriction
```